

**Table 3.11** Moore state transition table

Current State <i>S</i>	Input <i>A</i>	Next State <i>S'</i>
S0	0	S0
S0	1	S1
S1	0	S0
S1	1	S2
S2	0	S3
S2	1	S2
S3	0	S0
S3	1	S4
S4	0	S0
S4	1	S2

**Table 3.12** Moore output table

Current State <i>S</i>	Output <i>Y</i>
S0	0
S1	0
S2	0
S3	0
S4	1

**Table 3.13** Moore state transition table with state encodings

Current State			Input	Next State		
<i>S</i> <sub>2</sub>	<i>S</i> <sub>1</sub>	<i>S</i> <sub>0</sub>	<i>A</i>	<i>S'</i> <sub>2</sub>	<i>S'</i> <sub>1</sub>	<i>S'</i> <sub>0</sub>
0	0	0	0	0	0	0
0	0	0	1	0	0	1
0	0	1	0	0	0	0
0	0	1	1	0	1	0
0	1	0	0	0	1	1
0	1	0	1	0	1	0
0	1	1	0	0	0	0
0	1	1	1	1	0	0
1	0	0	0	0	0	0
1	0	0	1	0	1	0

**Table 3.14** Moore output table with state encodings

Current State <i>S</i> <sub>2</sub> <i>S</i> <sub>1</sub> <i>S</i> <sub>0</sub>	Output <i>Y</i>
0 0 0	0
0 0 1	0
0 1 0	0
0 1 1	0
1 0 0	1

**Table 3.15 Mealy state transition and output table**

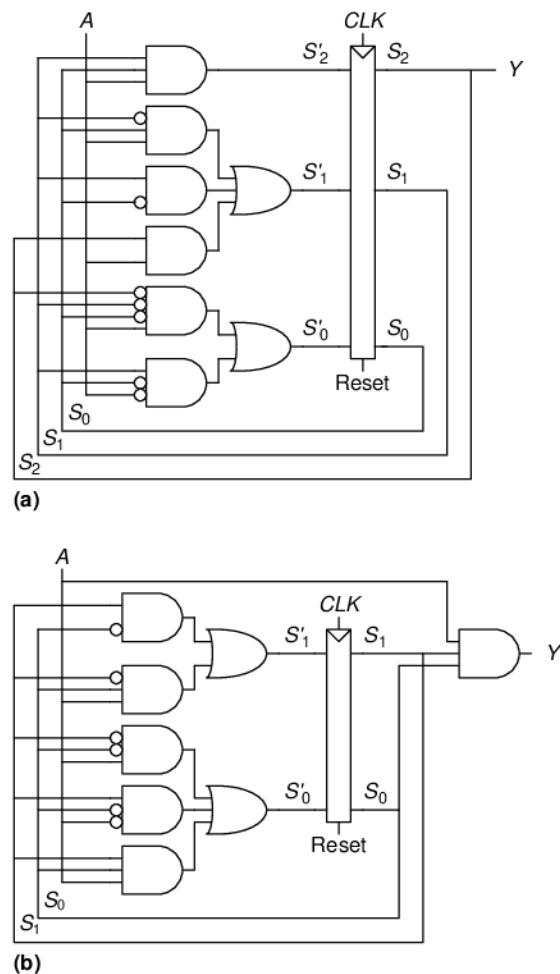
Current State S	Input A	Next State S'	Output Y
S0	0	S0	0
S0	1	S1	0
S1	0	S0	0
S1	1	S2	0
S2	0	S3	0
S2	1	S2	0
S3	0	S0	0
S3	1	S1	1

**Table 3.16 Mealy state transition and output table with state encodings**

Current State S <sub>1</sub> S <sub>0</sub>	Input A	Next State S' <sub>1</sub> S' <sub>0</sub>	Output Y
0 0	0	0 0	0
0 0	1	0 1	0
0 1	0	0 0	0
0 1	1	1 0	0
1 0	0	1 1	0
1 0	1	1 0	0
1 1	0	0 0	0
1 1	1	0 1	1

#### 3.4.4 Factoring State Machines

Designing complex FSMs is often easier if they can be broken down into multiple interacting simpler state machines such that the output of some machines is the input of others. This application of hierarchy and modularity is called *factoring* of state machines.



**Figure 3.31** FSM schematics for  
(a) Moore and (b) Mealy  
machines

---

#### Example 3.8 UNFACTORED AND FACTORED STATE MACHINES

Modify the traffic light controller from Section 3.4.1 to have a parade mode, which keeps the Bravado Boulevard light green while spectators and the band march to football games in scattered groups. The controller receives two more inputs:  $P$  and  $R$ . Asserting  $P$  for at least one cycle enters parade mode. Asserting  $R$  for at least one cycle leaves parade mode. When in parade mode, the controller proceeds through its usual sequence until  $L_B$  turns green, then remains in that state with  $L_B$  green until parade mode ends.

First, sketch a state transition diagram for a single FSM, as shown in Figure 3.33(a). Then, sketch the state transition diagrams for two interacting FSMs, as

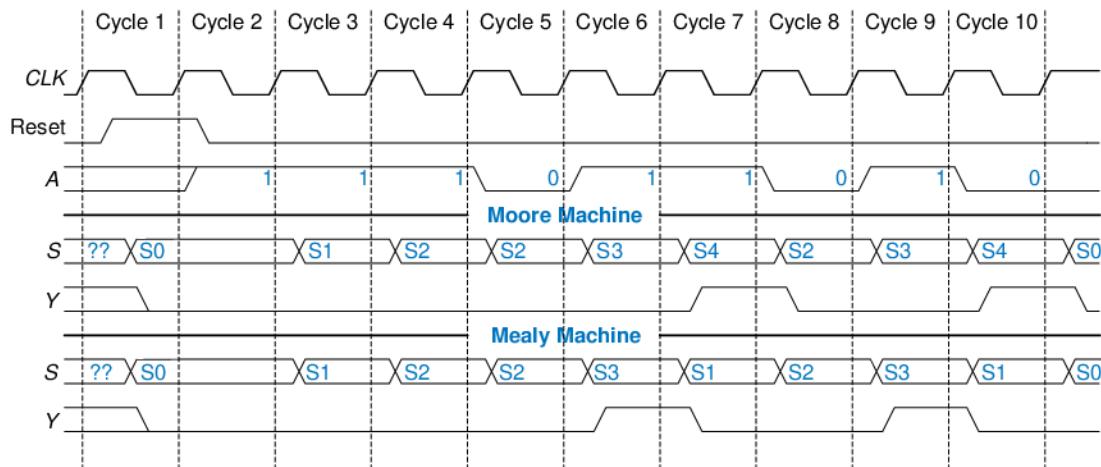


Figure 3.32 Timing diagrams for Moore and Mealy machines

shown in Figure 3.33(b). The Mode FSM asserts the output  $M$  when it is in parade mode. The Lights FSM controls the lights based on  $M$  and the traffic sensors,  $T_A$  and  $T_B$ .

**Solution:** Figure 3.34(a) shows the single FSM design. States S0 to S3 handle normal mode. States S4 to S7 handle parade mode. The two halves of the diagram are almost identical, but in parade mode, the FSM remains in S6 with a green light on Bravado Blvd. The  $P$  and  $R$  inputs control movement between these two halves. The FSM is messy and tedious to design. Figure 3.34(b) shows the factored FSM design. The mode FSM has two states to track whether the lights are in normal or parade mode. The Lights FSM is modified to remain in S2 while  $M$  is TRUE.

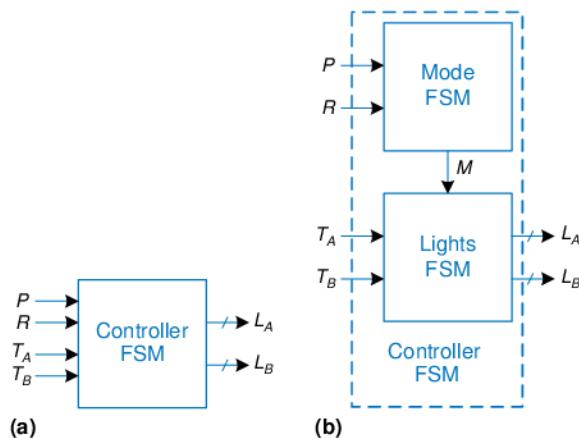
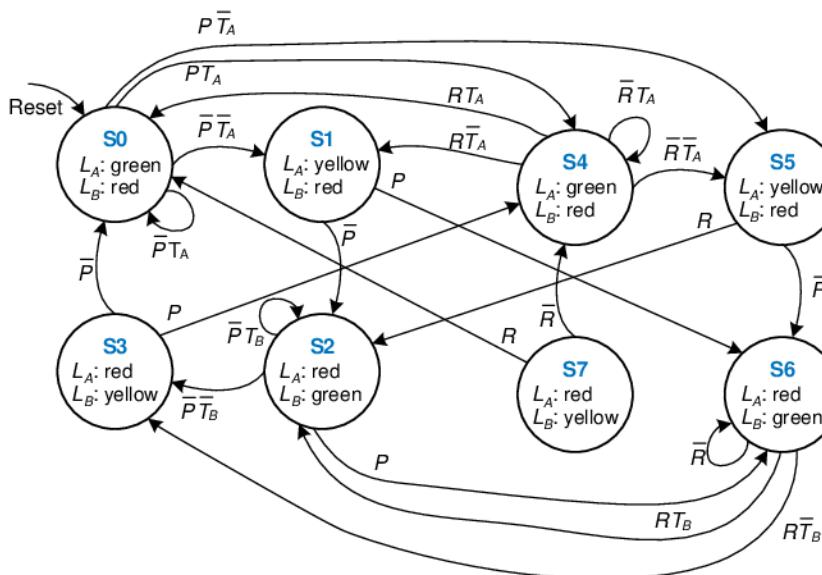
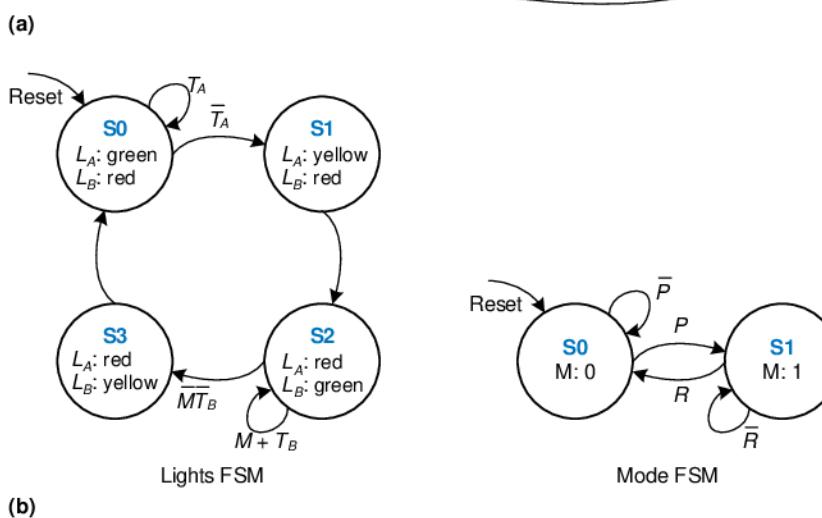


Figure 3.33 (a) single and (b) factored designs for modified traffic light controller FSM



**Figure 3.34** State transition diagrams: (a) unfactored, (b) factored



### 3.4.5 FSM Review

Finite state machines are a powerful way to systematically design sequential circuits from a written specification. Use the following procedure to design an FSM:

- ▶ Identify the inputs and outputs.
- ▶ Sketch a state transition diagram.

- ▶ For a Moore machine:
  - Write a state transition table.
  - Write an output table.
- ▶ For a Mealy machine:
  - Write a combined state transition and output table.
  - Select state encodings—your selection affects the hardware design.
  - Write Boolean equations for the next state and output logic.
  - Sketch the circuit schematic.

We will repeatedly use FSMs to design complex digital systems throughout this book.

### 3.5 TIMING OF SEQUENTIAL LOGIC

Recall that a flip-flop copies the input  $D$  to the output  $Q$  on the rising edge of the clock. This process is called *sampling D* on the clock edge. If  $D$  is *stable* at either 0 or 1 when the clock rises, this behavior is clearly defined. But what happens if  $D$  is changing at the same time the clock rises?

This problem is similar to that faced by a camera when snapping a picture. Imagine photographing a frog jumping from a lily pad into the lake. If you take the picture before the jump, you will see a frog on a lily pad. If you take the picture after the jump, you will see ripples in the water. But if you take it just as the frog jumps, you may see a blurred image of the frog stretching from the lily pad into the water. A camera is characterized by its *aperture time*, during which the object must remain still for a sharp image to be captured. Similarly, a sequential element has an aperture time around the clock edge, during which the input must be stable for the flip-flop to produce a well-defined output.

The aperture of a sequential element is defined by a *setup* time and a *hold* time, before and after the clock edge, respectively. Just as the static discipline limited us to using logic levels outside the forbidden zone, the *dynamic discipline* limits us to using signals that change outside the aperture time. By taking advantage of the dynamic discipline, we can think of time in discrete units called *clock cycles*, just as we think of signal levels as discrete 1's and 0's. A signal may glitch and oscillate wildly for some bounded amount of time. Under the dynamic discipline, we are concerned only about its final value at the end of the clock cycle, after it has settled to a stable value. Hence, we can simply write  $A[n]$ , the value of signal  $A$  at the end of the  $n^{\text{th}}$  clock cycle, where  $n$  is an integer, rather than  $A(t)$ , the value of  $A$  at some instant  $t$ , where  $t$  is any real number.

The clock period has to be long enough for all signals to settle. This sets a limit on the speed of the system. In real systems, the clock does not



reach all flip-flops at precisely the same time. This variation in time, called clock skew, further increases the necessary clock period.

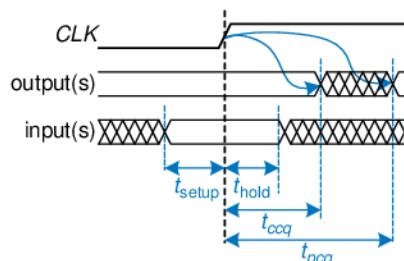
Sometimes it is impossible to satisfy the dynamic discipline, especially when interfacing with the real world. For example, consider a circuit with an input coming from a button. A monkey might press the button just as the clock rises. This can result in a phenomenon called metastability, where the flip-flop captures a value partway between 0 and 1 that can take an unlimited amount of time to resolve into a good logic value. The solution to such asynchronous inputs is to use a synchronizer, which has a very small (but nonzero) probability of producing an illegal logic value.

We expand on all of these ideas in the rest of this section.

### 3.5.1 The Dynamic Discipline

So far, we have focused on the functional specification of sequential circuits. Recall that a synchronous sequential circuit, such as a flip-flop or FSM, also has a timing specification, as illustrated in Figure 3.35. When the clock rises, the output (or outputs) may start to change after the clock-to-Q *contamination delay*,  $t_{ccq}$ , and must definitely settle to the final value within the clock-to-Q *propagation delay*,  $t_{pcq}$ . These represent the fastest and slowest delays through the circuit, respectively. For the circuit to sample its input correctly, the input (or inputs) must have stabilized at least some *setup time*,  $t_{\text{setup}}$ , before the rising edge of the clock and must remain stable for at least some *hold time*,  $t_{\text{hold}}$ , after the rising edge of the clock. The sum of the setup and hold times is called the *aperture time* of the circuit, because it is the total time for which the input must remain stable.

The *dynamic discipline* states that the inputs of a synchronous sequential circuit must be stable during the setup and hold aperture time around the clock edge. By imposing this requirement, we guarantee that the flip-flops sample signals while they are not changing. Because we are concerned only about the final values of the inputs at the time they are sampled, we can treat signals as discrete in time as well as in logic levels.



**Figure 3.35** Timing specification for synchronous sequential circuit

### 3.5.2 System Timing

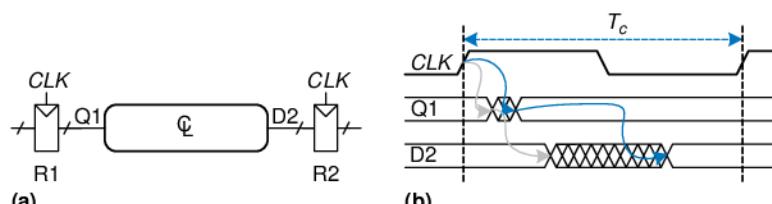
The *clock period* or *cycle time*,  $T_c$ , is the time between rising edges of a repetitive clock signal. Its reciprocal,  $f_c = 1/T_c$ , is the *clock frequency*. All else being the same, increasing the clock frequency increases the work that a digital system can accomplish per unit time. Frequency is measured in units of Hertz (Hz), or cycles per second: 1 megahertz (MHz) =  $10^6$  Hz, and 1 gigahertz (GHz) =  $10^9$  Hz.

Figure 3.36(a) illustrates a generic path in a synchronous sequential circuit whose clock period we wish to calculate. On the rising edge of the clock, register R1 produces output (or outputs) Q1. These signals enter a block of combinational logic, producing D2, the input (or inputs) to register R2. The timing diagram in Figure 3.36(b) shows that each output signal may start to change a contamination delay after its input change and settles to the final value within a propagation delay after its input settles. The gray arrows represent the contamination delay through R1 and the combinational logic, and the blue arrows represent the propagation delay through R1 and the combinational logic. We analyze the timing constraints with respect to the setup and hold time of the second register, R2.

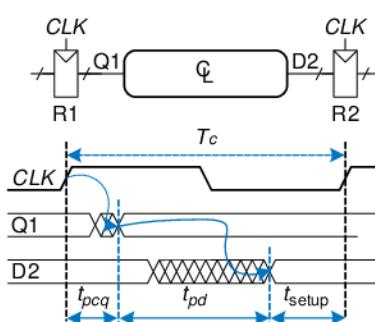
In the three decades from when one of the authors' families bought an Apple II+ computer to the present time of writing, microprocessor clock frequencies have increased from 1 MHz to several GHz, a factor of more than 1000. This speedup partially explains the revolutionary changes computers have made in society.

#### Setup Time Constraint

Figure 3.37 is the timing diagram showing only the maximum delay through the path, indicated by the blue arrows. To satisfy the setup time of R2, D2 must settle no later than the setup time before the next clock edge.



**Figure 3.36 Path between registers and timing diagram**



**Figure 3.37 Maximum delay for setup time constraint**

Hence, we find an equation for the minimum clock period:

$$T_c \geq t_{pcq} + t_{pd} + t_{\text{setup}} \quad (3.12)$$

In commercial designs, the clock period is often dictated by the Director of Engineering or by the marketing department (to ensure a competitive product). Moreover, the flip-flop clock-to- $Q$  propagation delay and setup time,  $t_{pcq}$  and  $t_{\text{setup}}$ , are specified by the manufacturer. Hence, we rearrange Equation 3.12 to solve for the maximum propagation delay through the combinational logic, which is usually the only variable under the control of the individual designer.

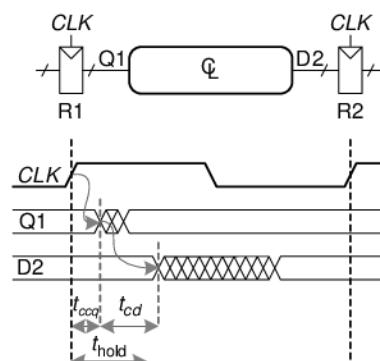
$$t_{pd} \leq T_c - (t_{pcq} + t_{\text{setup}}) \quad (3.13)$$

The term in parentheses,  $t_{pcq} + t_{\text{setup}}$ , is called the *sequencing overhead*. Ideally, the entire cycle time,  $T_c$ , would be available for useful computation in the combinational logic,  $t_{pd}$ . However, the sequencing overhead of the flip-flop cuts into this time. Equation 3.13 is called the *setup time constraint* or *max-delay constraint*, because it depends on the setup time and limits the maximum delay through combinational logic.

If the propagation delay through the combinational logic is too great,  $D_2$  may not have settled to its final value by the time  $R_2$  needs it to be stable and samples it. Hence,  $R_2$  may sample an incorrect result or even an illegal logic level, a level in the forbidden region. In such a case, the circuit will malfunction. The problem can be solved by increasing the clock period or by redesigning the combinational logic to have a shorter propagation delay.

#### Hold Time Constraint

The register  $R_2$  in Figure 3.36(a) also has a *hold time constraint*. Its input,  $D_2$ , must not change until some time,  $t_{\text{hold}}$ , after the rising edge of the clock. According to Figure 3.38,  $D_2$  might change as soon as  $t_{ccq} + t_{cd}$  after the rising edge of the clock.



**Figure 3.38** Minimum delay for hold time constraint

Hence, we find

$$t_{ccq} + t_{cd} \geq t_{hold} \quad (3.14)$$

Again,  $t_{ccq}$  and  $t_{hold}$  are characteristics of the flip-flop that are usually outside the designer's control. Rearranging, we can solve for the minimum contamination delay through the combinational logic:

$$t_{cd} \geq t_{hold} - t_{ccq} \quad (3.15)$$

Equation 3.15 is also called the *min-delay constraint* because it limits the minimum delay through combinational logic.

We have assumed that any logic elements can be connected to each other without introducing timing problems. In particular, we would expect that two flip-flops may be directly cascaded as in Figure 3.39 without causing hold time problems.

In such a case,  $t_{cd} = 0$  because there is no combinational logic between flip-flops. Substituting into Equation 3.15 yields the requirement that

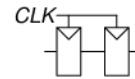
$$t_{hold} \leq t_{ccq} \quad (3.16)$$

In other words, a reliable flip-flop must have a hold time shorter than its contamination delay. Often, flip-flops are designed with  $t_{hold} = 0$ , so that Equation 3.16 is always satisfied. Unless noted otherwise, we will usually make that assumption and ignore the hold time constraint in this book.

Nevertheless, hold time constraints are critically important. If they are violated, the only solution is to increase the contamination delay through the logic, which requires redesigning the circuit. Unlike setup time constraints, they cannot be fixed by adjusting the clock period. Redesigning an integrated circuit and manufacturing the corrected design takes months and millions of dollars in today's advanced technologies, so *hold time violations* must be taken extremely seriously.

### Putting It All Together

Sequential circuits have setup and hold time constraints that dictate the maximum and minimum delays of the combinational logic between flip-flops. Modern flip-flops are usually designed so that the minimum delay through the combinational logic is 0—that is, flip-flops can be placed back-to-back. The maximum delay constraint limits the number of consecutive gates on the critical path of a high-speed circuit, because a high clock frequency means a short clock period.

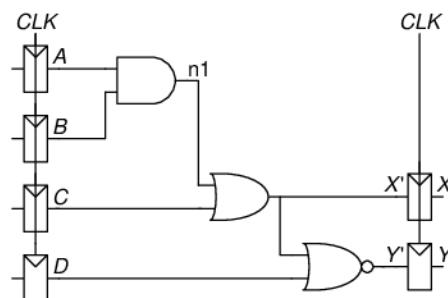


**Figure 3.39** Back-to-back flip-flops

---

### Example 3.9 TIMING ANALYSIS

Ben Bitdiddle designed the circuit in Figure 3.40. According to the data sheets for the components he is using, flip-flops have a clock-to-Q contamination delay



**Figure 3.40** Sample circuit for timing analysis

of 30 ps and a propagation delay of 80 ps. They have a setup time of 50 ps and a hold time of 60 ps. Each logic gate has a propagation delay of 40 ps and a contamination delay of 25 ps. Help Ben determine the maximum clock frequency and whether any hold time violations could occur. This process is called *timing analysis*.

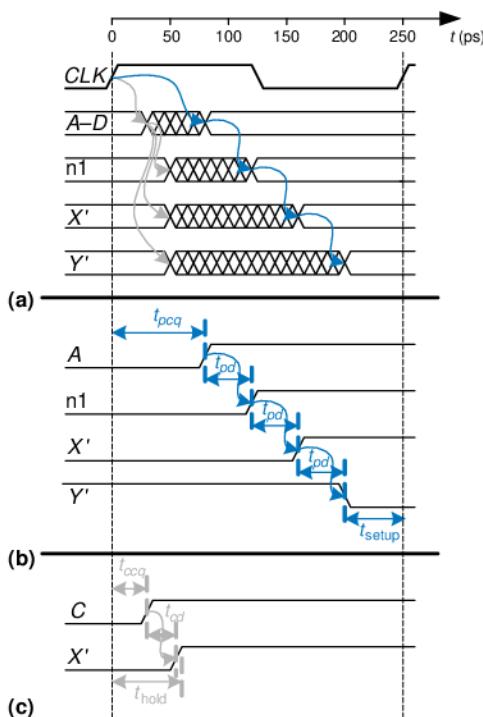
**Solution:** Figure 3.41(a) shows waveforms illustrating when the signals might change. The inputs,  $A$  to  $D$ , are registered, so they change shortly after  $CLK$  rises only.

The critical path occurs when  $B = 1$ ,  $C = 0$ ,  $D = 0$ , and  $A$  rises from 0 to 1, triggering  $n1$  to rise,  $X'$  to rise and  $Y'$  to fall, as shown in Figure 3.41(b). This path involves three gate delays. For the critical path, we assume that each gate requires its full propagation delay.  $Y'$  must setup before the next rising edge of the  $CLK$ . Hence, the minimum cycle time is

$$T_c \geq t_{pcq} + 3t_{pd} + t_{\text{setup}} = 80 + 3 \times 40 + 50 = 250 \text{ ps} \quad (3.17)$$

The maximum clock frequency is  $f_c = 1/T_c = 4 \text{ GHz}$ .

A short path occurs when  $A = 0$  and  $C$  rises, causing  $X'$  to rise, as shown in Figure 3.41(c). For the short path, we assume that each gate switches after only a contamination delay. This path involves only one gate delay, so it may occur after  $t_{ccq} + t_{cd} = 30 + 25 = 55 \text{ ps}$ . But recall that the flip-flop has a hold time of 60 ps, meaning that  $X'$  must remain stable for 60 ps after the rising edge of  $CLK$  for the flip-flop to reliably sample its value. In this case,  $X' = 0$  at the first rising edge of  $CLK$ , so we want the flip-flop to capture  $X = 0$ . Because  $X'$  did not hold stable long enough, the actual value of  $X$  is unpredictable. The circuit has a hold time violation and may behave erratically at any clock frequency.



**Figure 3.41** Timing diagram:  
(a) general case, (b) critical  
path, (c) short path

---

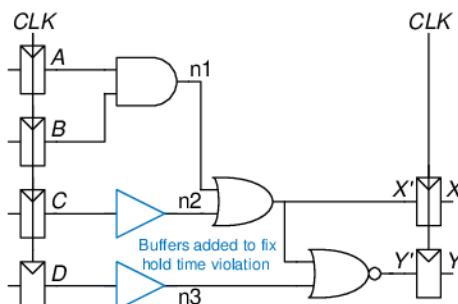
#### Example 3.10 FIXING HOLD TIME VIOLATIONS

Alyssa P. Hacker proposes to fix Ben's circuit by adding buffers to slow down the short paths, as shown in Figure 3.42. The buffers have the same delays as other gates. Determine the maximum clock frequency and whether any hold time problems could occur.

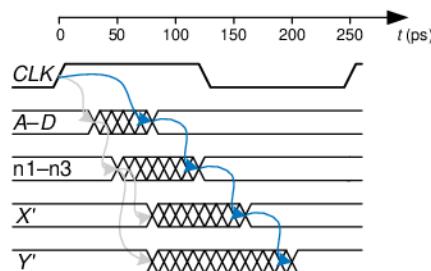
**Solution:** Figure 3.43 shows waveforms illustrating when the signals might change. The critical path from  $A$  to  $Y$  is unaffected, because it does not pass through any buffers. Therefore, the maximum clock frequency is still 4 GHz. However, the short paths are slowed by the contamination delay of the buffer. Now  $X'$  will not change until  $t_{ccq} + 2t_{cd} = 30 + 2 \times 25 = 80$  ps. This is after the 60 ps hold time has elapsed, so the circuit now operates correctly.

This example had an unusually long hold time to illustrate the point of hold time problems. Most flip-flops are designed with  $t_{hold} < t_{ccq}$  to avoid such problems. However, several high-performance microprocessors, including the Pentium 4, use an element called a *pulsed latch* in place of a flip-flop. The pulsed latch behaves like a flip-flop but has a short clock-to-Q delay and a long hold time. In general, adding buffers can usually, but not always, solve hold time problems without slowing the critical path.

---



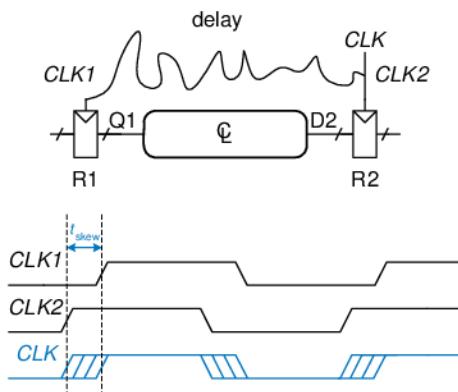
**Figure 3.42** Corrected circuit to fix hold time problem



**Figure 3.43** Timing diagram with buffers to fix hold time problem

### 3.5.3 Clock Skew\*

In the previous analysis, we assumed that the clock reaches all registers at exactly the same time. In reality, there is some variation in this time. This variation in clock edges is called *clock skew*. For example, the wires from the clock source to different registers may be of different lengths, resulting in slightly different delays, as shown in Figure 3.44. Noise also results in different delays. Clock gating, described in Section 3.2.5, further delays the clock. If some clocks are gated and others are not, there will be substantial skew between the gated and ungated clocks. In



**Figure 3.44** Clock skew caused by wire delay

Figure 3.44,  $CLK_2$  is *early* with respect to  $CLK_1$ , because the clock wire between the two registers follows a scenic route. If the clock had been routed differently,  $CLK_1$  might have been early instead. When doing timing analysis, we consider the worst-case scenario, so that we can guarantee that the circuit will work under all circumstances.

Figure 3.45 adds skew to the timing diagram from Figure 3.36. The heavy clock line indicates the latest time at which the clock signal might reach any register; the hashed lines show that the clock might arrive up to  $t_{\text{skew}}$  earlier.

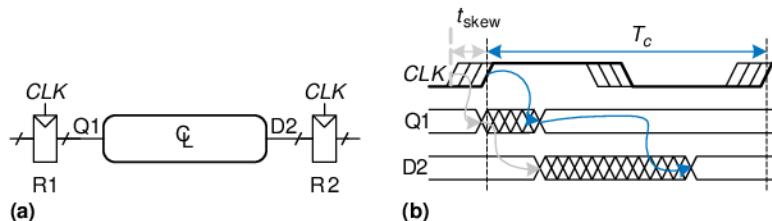
First, consider the setup time constraint shown in Figure 3.46. In the worst case, R1 receives the latest skewed clock and R2 receives the earliest skewed clock, leaving as little time as possible for data to propagate between the registers.

The data propagates through the register and combinational logic and must setup before R2 samples it. Hence, we conclude that

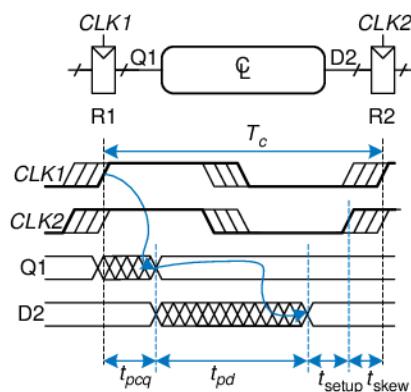
$$T_c \geq t_{\text{pcq}} + t_{\text{pd}} + t_{\text{setup}} + t_{\text{skew}} \quad (3.18)$$

$$t_{\text{pd}} \leq T_c - (t_{\text{pcq}} + t_{\text{setup}} + t_{\text{skew}}) \quad (3.19)$$

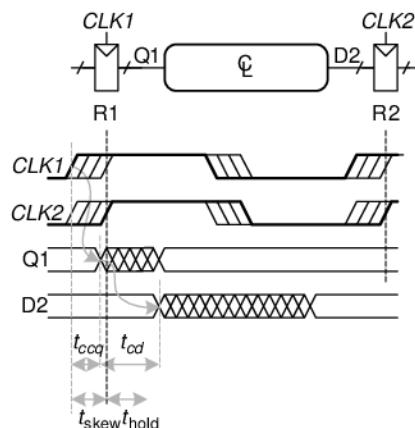
Next, consider the hold time constraint shown in Figure 3.47. In the worst case, R1 receives an early skewed clock,  $CLK_1$ , and R2 receives a late skewed clock,  $CLK_2$ . The data zips through the register



**Figure 3.45** Timing diagram with clock skew



**Figure 3.46** Setup time constraint with clock skew



**Figure 3.47** Hold time constraint with clock skew

and combinational logic but must not arrive until a hold time after the late clock. Thus, we find that

$$t_{ccq} + t_{cd} \geq t_{hold} + t_{skew} \quad (3.20)$$

$$t_{cd} \geq t_{hold} + t_{skew} - t_{ccq} \quad (3.21)$$

In summary, clock skew effectively increases both the setup time and the hold time. It adds to the sequencing overhead, reducing the time available for useful work in the combinational logic. It also increases the required minimum delay through the combinational logic. Even if  $t_{hold} = 0$ , a pair of back-to-back flip-flops will violate Equation 3.21 if  $t_{skew} > t_{ccq}$ . To prevent serious hold time failures, designers must not permit too much clock skew. Sometimes flip-flops are intentionally designed to be particularly slow (i.e., large  $t_{ccq}$ ), to prevent hold time problems even when the clock skew is substantial.

---

#### Example 3.11 TIMING ANALYSIS WITH CLOCK SKEW

Revisit Example 3.9 and assume that the system has 50 ps of clock skew.

**Solution:** The critical path remains the same, but the setup time is effectively increased by the skew. Hence, the minimum cycle time is

$$\begin{aligned} T_c &\geq t_{pcq} + 3t_{pd} + t_{setup} + t_{skew} \\ &= 80 + 3 \times 40 + 50 + 50 = 300 \text{ ps} \end{aligned} \quad (3.22)$$

The maximum clock frequency is  $f_c = 1/T_c = 3.33 \text{ GHz}$ .

The short path also remains the same at 55 ps. The hold time is effectively increased by the skew to  $60 + 50 = 110 \text{ ps}$ , which is much greater than 55 ps. Hence, the circuit will violate the hold time and malfunction at any frequency. The circuit violated the hold time constraint even without skew. Skew in the system just makes the violation worse.

---

---

**Example 3.12 FIXING HOLD TIME VIOLATIONS**

Revisit Example 3.10 and assume that the system has 50 ps of clock skew.

**Solution:** The critical path is unaffected, so the maximum clock frequency remains 3.33 GHz.

The short path increases to 80 ps. This is still less than  $t_{hold} + t_{skew} = 110$  ps, so the circuit still violates its hold time constraint.

To fix the problem, even more buffers could be inserted. Buffers would need to be added on the critical path as well, reducing the clock frequency. Alternatively, a better flip-flop with a shorter hold time might be used.

---

### 3.5.4 Metastability

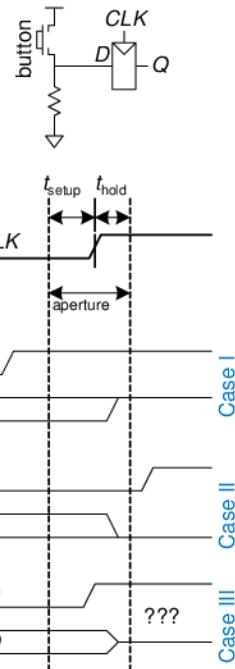
As noted earlier, it is not always possible to guarantee that the input to a sequential circuit is stable during the aperture time, especially when the input arrives from the external world. Consider a button connected to the input of a flip-flop, as shown in Figure 3.48. When the button is not pressed,  $D = 0$ . When the button is pressed,  $D = 1$ . A monkey presses the button at some random time relative to the rising edge of  $CLK$ . We want to know the output  $Q$  after the rising edge of  $CLK$ . In Case I, when the button is pressed much before  $CLK$ ,  $Q = 1$ .

In Case II, when the button is not pressed until long after  $CLK$ ,  $Q = 0$ . But in Case III, when the button is pressed sometime between  $t_{setup}$  before  $CLK$  and  $t_{hold}$  after  $CLK$ , the input violates the dynamic discipline and the output is undefined.

#### Metastable State

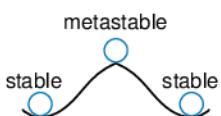
In reality, when a flip-flop samples an input that is changing during its aperture, the output  $Q$  may momentarily take on a voltage between 0 and  $V_{DD}$  that is in the forbidden zone. This is called a *metastable* state. Eventually, the flip-flop will resolve the output to a *stable state* of either 0 or 1. However, the *resolution time* required to reach the stable state is unbounded.

The metastable state of a flip-flop is analogous to a ball on the summit of a hill between two valleys, as shown in Figure 3.49. The two valleys are stable states, because a ball in the valley will remain there as long as it is not disturbed. The top of the hill is called metastable because the ball would remain there if it were perfectly balanced. But because nothing is perfect, the ball will eventually roll to one side or the other. The time required for this change to occur depends on how nearly well balanced the ball originally was. Every bistable device has a metastable state between the two stable states.



**Figure 3.48** Input changing before, after, or during aperture





**Figure 3.49** Stable and metastable states

### Resolution Time

If a flip-flop input changes at a random time during the clock cycle, the resolution time,  $t_{res}$ , required to resolve to a stable state is also a random variable. If the input changes outside the aperture, then  $t_{res} = t_{pcq}$ . But if the input happens to change within the aperture,  $t_{res}$  can be substantially longer. Theoretical and experimental analyses (see Section 3.5.6) have shown that the probability that the resolution time,  $t_{res}$ , exceeds some arbitrary time,  $t$ , decreases exponentially with  $t$ :

$$P(t_{res} > t) = \frac{T_0}{T_c} e^{-\frac{t}{\tau}} \quad (3.23)$$

where  $T_c$  is the clock period, and  $T_0$  and  $\tau$  are characteristic of the flip-flop. The equation is valid only for  $t$  substantially longer than  $t_{pcq}$ .

Intuitively,  $T_0/T_c$  describes the probability that the input changes at a bad time (i.e., during the aperture time); this probability decreases with the cycle time,  $T_c$ .  $\tau$  is a time constant indicating how fast the flip-flop moves away from the metastable state; it is related to the delay through the cross-coupled gates in the flip-flop.

In summary, if the input to a bistable device such as a flip-flop changes during the aperture time, the output may take on a metastable value for some time before resolving to a stable 0 or 1. The amount of time required to resolve is unbounded, because for any finite time,  $t$ , the probability that the flip-flop is still metastable is nonzero. However, this probability drops off exponentially as  $t$  increases. Therefore, if we wait long enough, much longer than  $t_{pcq}$ , we can expect with exceedingly high probability that the flip-flop will reach a valid logic level.

### 3.5.5 Synchronizers

Asynchronous inputs to digital systems from the real world are inevitable. Human input is asynchronous, for example. If handled carelessly, these asynchronous inputs can lead to metastable voltages within the system, causing erratic system failures that are extremely difficult to track down and correct. The goal of a digital system designer should be to ensure that, given asynchronous inputs, the probability of encountering a metastable voltage is sufficiently small. “Sufficiently” depends on the context. For a digital cell phone, perhaps one failure in 10 years is acceptable, because the user can always turn the phone off and back on if it locks up. For a medical device, one failure in the expected life of the universe ( $10^{10}$  years) is a better target. To guarantee good logic levels, all asynchronous inputs should be passed through *synchronizers*.

A synchronizer, shown in Figure 3.50, is a device that receives an asynchronous input,  $D$ , and a clock,  $CLK$ . It produces an output,  $Q$ , within a bounded amount of time; the output has a valid logic level with extremely high probability. If  $D$  is stable during the aperture,  $Q$  should take on the same value as  $D$ . If  $D$  changes during the aperture,  $Q$  may take on either a HIGH or LOW value but must not be metastable.

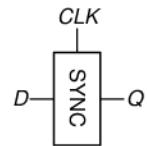
Figure 3.51 shows a simple way to build a synchronizer out of two flip-flops. F1 samples  $D$  on the rising edge of  $CLK$ . If  $D$  is changing at that time, the output D2 may be momentarily metastable. If the clock period is long enough, D2 will, with high probability, resolve to a valid logic level before the end of the period. F2 then samples D2, which is now stable, producing a good output  $Q$ .

We say that a synchronizer *fails* if  $Q$ , the output of the synchronizer, becomes metastable. This may happen if D2 has not resolved to a valid level by the time it must setup at F2—that is, if  $t_{res} > T_c - t_{setup}$ . According to Equation 3.23, the probability of failure for a single input change at a random time is

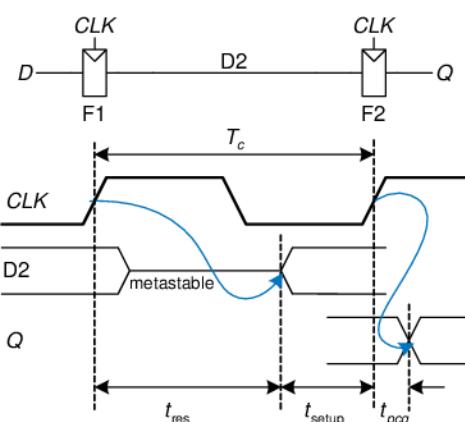
$$P(\text{failure}) = \frac{T_0}{T_c} e^{-\frac{T_c - t_{\text{setup}}}{\tau}} \quad (3.24)$$

The probability of failure,  $P(\text{failure})$ , is the probability that the output,  $Q$ , will be metastable upon a single change in  $D$ . If  $D$  changes once per second, the probability of failure per second is just  $P(\text{failure})$ . However, if  $D$  changes  $N$  times per second, the probability of failure per second is  $N$  times as great:

$$P(\text{failure})/\text{sec} = N \frac{T_0}{T_c} e^{-\frac{T_c - t_{\text{setup}}}{\tau}} \quad (3.25)$$



**Figure 3.50 Synchronizer symbol**



**Figure 3.51 Simple synchronizer**

System reliability is usually measured in *mean time between failures* (MTBF). As the name might suggest, MTBF is the average amount of time between failures of the system. It is the reciprocal of the probability that the system will fail in any given second

$$MTBF = \frac{1}{P(\text{failure})/\text{sec}} = \frac{T_c e^{\frac{T_c - t_{\text{setup}}}{\tau}}}{NT_0} \quad (3.26)$$

Equation 3.26 shows that the MTBF improves exponentially as the synchronizer waits for a longer time,  $T_c$ . For most systems, a synchronizer that waits for one clock cycle provides a safe MTBF. In exceptionally high-speed systems, waiting for more cycles may be necessary.

#### **Example 3.13 SYNCHRONIZER FOR FSM INPUT**

The traffic light controller FSM from Section 3.4.1 receives asynchronous inputs from the traffic sensors. Suppose that a synchronizer is used to guarantee stable inputs to the controller. Traffic arrives on average 0.2 times per second. The flip-flops in the synchronizer have the following characteristics:  $\tau = 200$  ps,  $T_0 = 150$  ps, and  $t_{\text{setup}} = 500$  ps. How long must the synchronizer clock period be for the MTBF to exceed 1 year?

**Solution:** 1 year  $\approx \pi \times 10^7$  seconds. Solve Equation 3.26.

$$\pi \times 10^7 = \frac{T_c e^{\frac{T_c - 500 \times 10^{-12}}{200 \times 10^{-12}}}}{(0.2)(150 \times 10^{-12})} \quad (3.27)$$

This equation has no closed form solution. However, it is easy enough to solve by guess and check. In a spreadsheet, try a few values of  $T_c$  and calculate the MTBF until discovering the value of  $T_c$  that gives an MTBF of 1 year:  $T_c = 3.036$  ns.

#### **3.5.6 Derivation of Resolution Time\***

Equation 3.23 can be derived using a basic knowledge of circuit theory, differential equations, and probability. This section can be skipped if you are not interested in the derivation or if you are unfamiliar with the mathematics.

A flip-flop output will be metastable after some time,  $t$ , if the flip-flop samples a changing input (causing a metastable condition) and the output does not resolve to a valid level within that time after the clock edge. Symbolically, this can be expressed as

$$P(t_{\text{res}} > t) = P(\text{samples changing input}) \times P(\text{unresolved}) \quad (3.28)$$

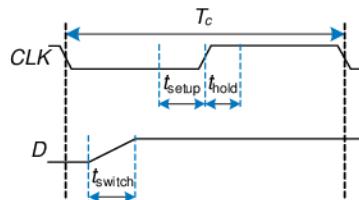


Figure 3.52 Input timing

We consider each probability term individually. The asynchronous input signal switches between 0 and 1 in some time,  $t_{\text{switch}}$ , as shown in Figure 3.52. The probability that the input changes during the aperture around the clock edge is

$$P(\text{samples changing input}) = \frac{t_{\text{switch}} + t_{\text{setup}} + t_{\text{hold}}}{T_c} \quad (3.29)$$

If the flip-flop does enter metastability—that is, with probability  $P(\text{samples changing input})$ —the time to resolve from metastability depends on the inner workings of the circuit. This resolution time determines  $P(\text{unresolved})$ , the probability that the flip-flop has not yet resolved to a valid logic level after a time  $t$ . The remainder of this section analyzes a simple model of a bistable device to estimate this probability.

A bistable device uses storage with positive feedback. Figure 3.53(a) shows this feedback implemented with a pair of inverters; this circuit's behavior is representative of most bistable elements. A pair of inverters behaves like a buffer. Let us model it as having the symmetric DC transfer characteristics shown in Figure 3.53(b), with a slope of  $G$ . The buffer can deliver only a finite amount of output current; we can model this as an output resistance,  $R$ . All real circuits also have some capacitance,  $C$ , that must be charged up. Charging the capacitor through the resistor causes an RC delay, preventing the buffer from switching instantaneously. Hence, the complete circuit model is shown in Figure 3.53(c), where  $v_{\text{out}}(t)$  is the voltage of interest conveying the state of the bistable device.

The metastable point for this circuit is  $v_{\text{out}}(t) = v_{\text{in}}(t) = V_{DD}/2$ ; if the circuit began at exactly that point, it would remain there indefinitely in the

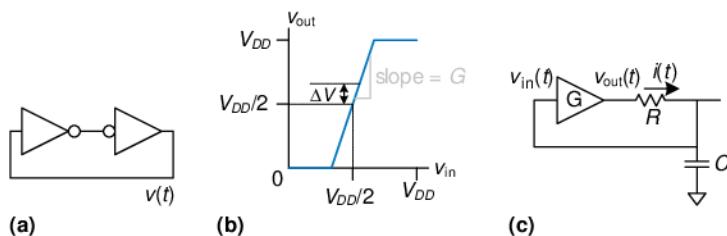


Figure 3.53 Circuit model of bistable device

absence of noise. Because voltages are continuous variables, the chance that the circuit will begin at exactly the metastable point is vanishingly small. However, the circuit might begin at time 0 near metastability at  $v_{\text{out}}(0) = V_{DD}/2 + \Delta V$  for some small offset  $\Delta V$ . In such a case, the positive feedback will eventually drive  $v_{\text{out}}(t)$  to  $V_{DD}$  if  $\Delta V > 0$  and to 0 if  $\Delta V < 0$ . The time required to reach  $V_{DD}$  or 0 is the resolution time of the bistable device.

The DC transfer characteristic is nonlinear, but it appears linear near the metastable point, which is the region of interest to us. Specifically, if  $v_{\text{in}}(t) = V_{DD}/2 + \Delta V/G$ , then  $v_{\text{out}}(t) = V_{DD}/2 + \Delta V$  for small  $\Delta V$ . The current through the resistor is  $i(t) = (v_{\text{out}}(t) - v_{\text{in}}(t))/R$ . The capacitor charges at a rate  $dv_{\text{in}}(t)/dt = i(t)/C$ . Putting these facts together, we find the governing equation for the output voltage.

$$\frac{dv_{\text{out}}(t)}{dt} = \frac{(G-1)}{RC} \left[ v_{\text{out}}(t) - \frac{V_{DD}}{2} \right] \quad (3.30)$$

This is a linear first-order differential equation. Solving it with the initial condition  $v_{\text{out}}(0) = V_{DD}/2 + \Delta V$  gives

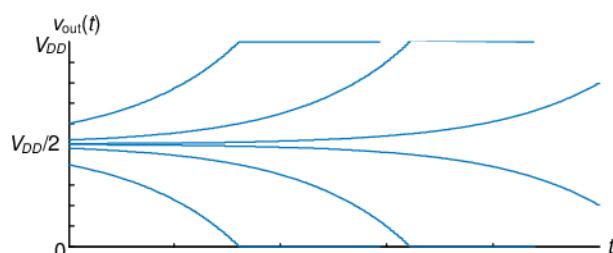
$$v_{\text{out}}(t) = \frac{V_{DD}}{2} + \Delta V e^{\frac{(G-1)t}{RC}} \quad (3.31)$$

Figure 3.54 plots trajectories for  $v_{\text{out}}(t)$  given various starting points.  $v_{\text{out}}(t)$  moves exponentially away from the metastable point  $V_{DD}/2$  until it saturates at  $V_{DD}$  or 0. The output voltage eventually resolves to 1 or 0. The amount of time this takes depends on the initial voltage offset ( $\Delta V$ ) from the metastable point ( $V_{DD}/2$ ).

Solving Equation 3.31 for the resolution time  $t_{\text{res}}$ , such that  $v_{\text{out}}(t_{\text{res}}) = V_{DD}$  or 0, gives

$$|\Delta V| e^{\frac{(G-1)t_{\text{res}}}{RC}} = \frac{V_{DD}}{2} \quad (3.32)$$

$$t_{\text{res}} = \frac{RC}{G-1} \ln \frac{V_{DD}}{2|\Delta V|} \quad (3.33)$$



**Figure 3.54** Resolution trajectories

In summary, the resolution time increases if the bistable device has high resistance or capacitance that causes the output to change slowly. It decreases if the bistable device has high *gain*,  $G$ . The resolution time also increases logarithmically as the circuit starts closer to the metastable point ( $\Delta V \rightarrow 0$ ).

Define  $\tau$  as  $\frac{RC}{G-1}$ . Solving Equation 3.33 for  $\Delta V$  finds the initial offset,  $\Delta V_{res}$ , that gives a particular resolution time,  $t_{res}$ :

$$\Delta V_{res} = \frac{V_{DD}}{2} e^{-t_{res}/\tau} \quad (3.34)$$

Suppose that the bistable device samples the input while it is changing. It measures a voltage,  $v_{in}(0)$ , which we will assume is uniformly distributed between 0 and  $V_{DD}$ . The probability that the output has not resolved to a legal value after time  $t_{res}$  depends on the probability that the initial offset is sufficiently small. Specifically, the initial offset on  $v_{out}$  must be less than  $\Delta V_{res}$ , so the initial offset on  $v_{in}$  must be less than  $\Delta V_{res}/G$ . Then the probability that the bistable device samples the input at a time to obtain a sufficiently small initial offset is

$$P(\text{unresolved}) = P\left(v_{in}(0) - \frac{V_{DD}}{2} < \frac{\Delta V_{res}}{G}\right) = \frac{2\Delta V_{res}}{GV_{DD}} \quad (3.35)$$

Putting this all together, the probability that the resolution time exceeds some time,  $t$ , is given by the following equation:

$$P(t_{res} > t) = \frac{t_{\text{switch}} + t_{\text{setup}} + t_{\text{hold}}}{GT_c} e^{-\frac{t}{\tau}} \quad (3.36)$$

Observe that Equation 3.36 is in the form of Equation 3.23, where  $T_0 = (t_{\text{switch}} + t_{\text{setup}} + t_{\text{hold}})/G$  and  $\tau = RC/(G-1)$ . In summary, we have derived Equation 3.23 and shown how  $T_0$  and  $\tau$  depend on physical properties of the bistable device.

### 3.6 PARALLELISM

The speed of a system is measured in latency and throughput of tokens moving through a system. We define a *token* to be a group of inputs that are processed to produce a group of outputs. The term conjures up the notion of placing subway tokens on a circuit diagram and moving them around to visualize data moving through the circuit. The *latency* of a system is the time required for one token to pass through the system from start to end. The *throughput* is the number of tokens that can be produced per unit time.



---

**Example 3.14 COOKIE THROUGHPUT AND LATENCY**

Ben Bitdiddle is throwing a milk and cookies party to celebrate the installation of his traffic light controller. It takes him 5 minutes to roll cookies and place them on his tray. It then takes 15 minutes for the cookies to bake in the oven. Once the cookies are baked, he starts another tray. What is Ben's throughput and latency for a tray of cookies?

**Solution:** In this example, a tray of cookies is a token. The latency is  $1/3$  hour per tray. The throughput is 3 trays/hour.

---

As you might imagine, the throughput can be improved by processing several tokens at the same time. This is called *parallelism*, and it comes in two forms: spatial and temporal. With *spatial parallelism*, multiple copies of the hardware are provided so that multiple tasks can be done at the same time. With *temporal parallelism*, a task is broken into stages, like an assembly line. Multiple tasks can be spread across the stages. Although each task must pass through all stages, a *different* task will be in each stage at any given time so multiple tasks can overlap. Temporal parallelism is commonly called *pipelining*. Spatial parallelism is sometimes just called parallelism, but we will avoid that naming convention because it is ambiguous.

---

**Example 3.15 COOKIE PARALLELISM**

Ben Bitdiddle has hundreds of friends coming to his party and needs to bake cookies faster. He is considering using spatial and/or temporal parallelism.

**Spatial Parallelism:** Ben asks Alyssa P. Hacker to help out. She has her own cookie tray and oven.

**Temporal Parallelism:** Ben gets a second cookie tray. Once he puts one cookie tray in the oven, he starts rolling cookies on the other tray rather than waiting for the first tray to bake.

What is the throughput and latency using spatial parallelism? Using temporal parallelism? Using both?

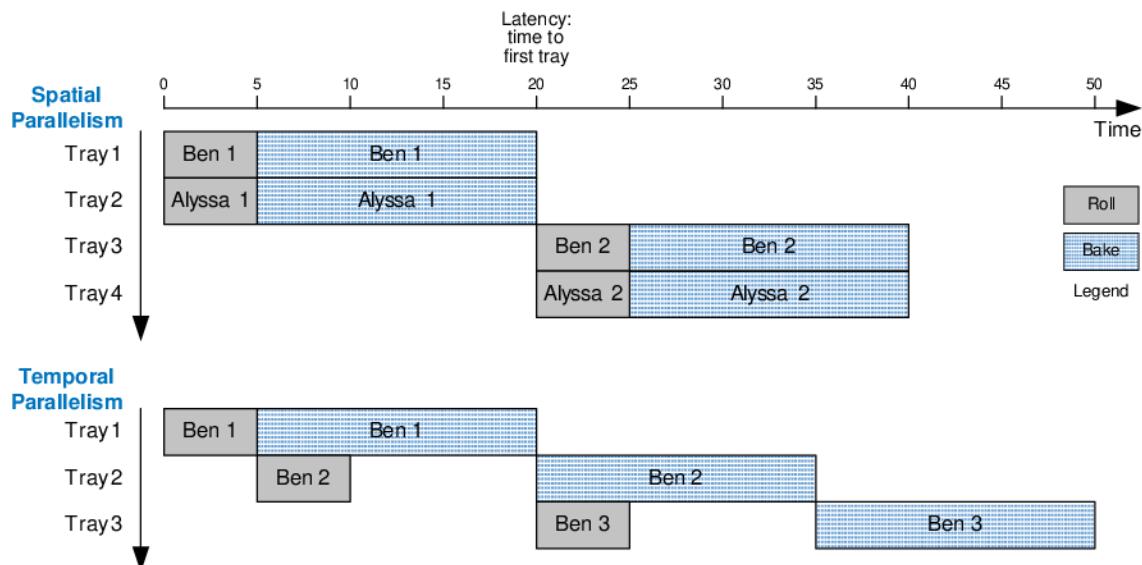
**Solution:** The latency is the time required to complete one task from start to finish. In all cases, the latency is  $1/3$  hour. If Ben starts with no cookies, the latency is the time needed for him to produce the first cookie tray.

The throughput is the number of cookie trays per hour. With spatial parallelism, Ben and Alyssa each complete one tray every 20 minutes. Hence, the throughput doubles, to 6 trays/hour. With temporal parallelism, Ben puts a new tray in the oven every 15 minutes, for a throughput of 4 trays/hour. These are illustrated in Figure 3.55.

---

If Ben and Alyssa use both techniques, they can bake 8 trays/hour.

---



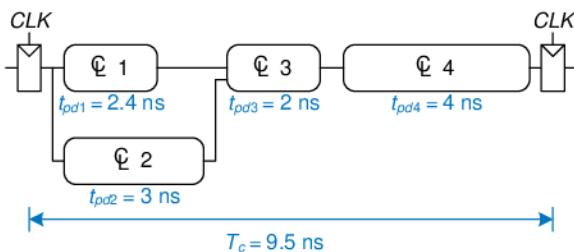
**Figure 3.55** Spatial and temporal parallelism in the cookie kitchen

Consider a task with latency  $L$ . In a system with no parallelism, the throughput is  $1/L$ . In a spatially parallel system with  $N$  copies of the hardware, the throughput is  $N/L$ . In a temporally parallel system, the task is ideally broken into  $N$  steps, or stages, of equal length. In such a case, the throughput is also  $N/L$ , and only one copy of the hardware is required. However, as the cookie example showed, finding  $N$  steps of equal length is often impractical. If the longest step has a latency  $L_1$ , the pipelined throughput is  $1/L_1$ .

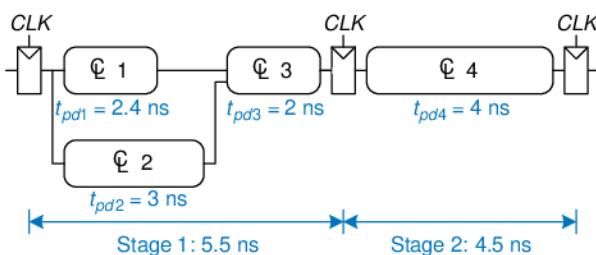
Pipelining (temporal parallelism) is particularly attractive because it speeds up a circuit without duplicating the hardware. Instead, registers are placed between blocks of combinational logic to divide the logic into shorter stages that can run with a faster clock. The registers prevent a token in one pipeline stage from catching up with and corrupting the token in the next stage.

Figure 3.56 shows an example of a circuit with no pipelining. It contains four blocks of logic between the registers. The critical path passes through blocks 2, 3, and 4. Assume that the register has a clock-to-Q propagation delay of 0.3 ns and a setup time of 0.2 ns. Then the cycle time is  $T_c = 0.3 + 3 + 2 + 4 + 0.2 = 9.5$  ns. The circuit has a latency of 9.5 ns and a throughput of  $1/9.5$  ns = 105 MHz.

Figure 3.57 shows the same circuit partitioned into a two-stage pipeline by adding a register between blocks 3 and 4. The first stage has a minimum clock period of  $0.3 + 3 + 2 + 0.2 = 5.5$  ns. The second



**Figure 3.56** Circuit with no pipelining

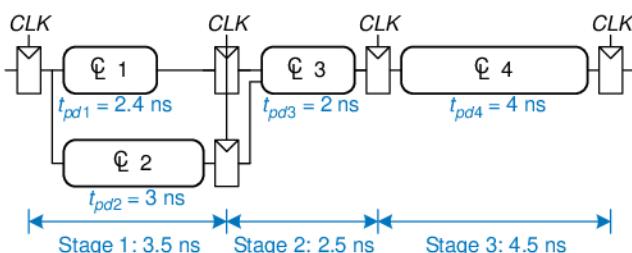


**Figure 3.57** Circuit with two-stage pipeline

stage has a minimum clock period of  $0.3 + 4 + 0.2 = 4.5 \text{ ns}$ . The clock must be slow enough for all stages to work. Hence,  $T_c = 5.5 \text{ ns}$ . The latency is two clock cycles, or 11 ns. The throughput is  $1/5.5 \text{ ns} = 182 \text{ MHz}$ . This example shows that, in a real circuit, pipelining with two stages almost doubles the throughput and slightly increases the latency. In comparison, ideal pipelining would exactly double the throughput at no penalty in latency. The discrepancy comes about because the circuit cannot be divided into two exactly equal halves and because the registers introduce more sequencing overhead.

Figure 3.58 shows the same circuit partitioned into a three-stage pipeline. Note that two more registers are needed to store the results of blocks 1 and 2 at the end of the first pipeline stage. The cycle time is now limited by the third stage to 4.5 ns. The latency is three cycles, or 13.5 ns. The throughput is  $1/4.5 \text{ ns} = 222 \text{ MHz}$ . Again, adding a pipeline stage improves throughput at the expense of some latency.

Although these techniques are powerful, they do not apply to all situations. The bane of parallelism is *dependencies*. If a current task is



**Figure 3.58** Circuit with three-stage pipeline

dependent on the result of a prior task, rather than just prior steps in the current task, the task cannot start until the prior task has completed. For example, if Ben wants to check that the first tray of cookies tastes good before he starts preparing the second, he has a dependency that prevents pipelining or parallel operation. Parallelism is one of the most important techniques for designing high-performance microprocessors. Chapter 7 discusses pipelining further and shows examples of handling dependencies.

### 3.7 SUMMARY

This chapter has described the analysis and design of sequential logic. In contrast to combinational logic, whose outputs depend only on the current inputs, sequential logic outputs depend on both current and prior inputs. In other words, sequential logic remembers information about prior inputs. This memory is called the state of the logic.

Sequential circuits can be difficult to analyze and are easy to design incorrectly, so we limit ourselves to a small set of carefully designed building blocks. The most important element for our purposes is the flip-flop, which receives a clock and an input,  $D$ , and produces an output,  $Q$ . The flip-flop copies  $D$  to  $Q$  on the rising edge of the clock and otherwise remembers the old state of  $Q$ . A group of flip-flops sharing a common clock is called a register. Flip-flops may also receive reset or enable control signals.

Although many forms of sequential logic exist, we discipline ourselves to use synchronous sequential circuits because they are easy to design. Synchronous sequential circuits consist of blocks of combinational logic separated by clocked registers. The state of the circuit is stored in the registers and updated only on clock edges.

Finite state machines are a powerful technique for designing sequential circuits. To design an FSM, first identify the inputs and outputs of the machine and sketch a state transition diagram, indicating the states and the transitions between them. Select an encoding for the states, and rewrite the diagram as a state transition table and output table, indicating the next state and output given the current state and input. From these tables, design the combinational logic to compute the next state and output, and sketch the circuit.

Synchronous sequential circuits have a timing specification including the clock-to- $Q$  propagation and contamination delays,  $t_{pcq}$  and  $t_{ccq}$ , and the setup and hold times,  $t_{\text{setup}}$  and  $t_{\text{hold}}$ . For correct operation, their inputs must be stable during an aperture time that starts a setup time before the rising edge of the clock and ends a hold time after the rising edge of the clock. The minimum cycle time,  $T_c$ , of the system is equal to the propagation delay,  $t_{pd}$ , through the combinational logic plus

Anyone who could invent logic whose outputs depend on future inputs would be fabulously wealthy!

$t_{pcq} + t_{\text{setup}}$  of the register. For correct operation, the contamination delay through the register and combinational logic must be greater than  $t_{\text{hold}}$ . Despite the common misconception to the contrary, hold time does not affect the cycle time.

Overall system performance is measured in latency and throughput. The latency is the time required for a token to pass from start to end. The throughput is the number of tokens that the system can process per unit time. Parallelism improves the system throughput.

## Exercises

**Exercise 3.1** Given the input waveforms shown in Figure 3.59, sketch the output,  $Q$ , of an SR latch.



Figure 3.59 Input waveform of SR latch

**Exercise 3.2** Given the input waveforms shown in Figure 3.60, sketch the output,  $Q$ , of a D latch.



Figure 3.60 Input waveform of D latch or flip-flop

**Exercise 3.3** Given the input waveforms shown in Figure 3.60, sketch the output,  $Q$ , of a D flip-flop.

**Exercise 3.4** Is the circuit in Figure 3.61 combinational logic or sequential logic? Explain in a simple fashion what the relationship is between the inputs and outputs. What would you call this circuit?

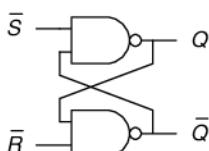


Figure 3.61 Mystery circuit

**Exercise 3.5** Is the circuit in Figure 3.62 combinational logic or sequential logic? Explain in a simple fashion what the relationship is between the inputs and outputs. What would you call this circuit?

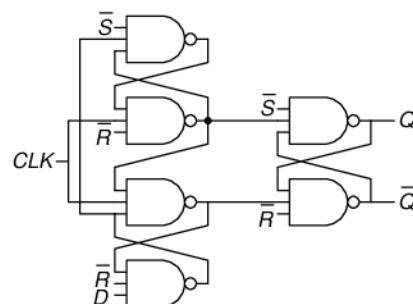


Figure 3.62 Mystery circuit

**Exercise 3.6** The *toggle (T) flip-flop* has one input, *CLK*, and one output, *Q*. On each rising edge of *CLK*, *Q* toggles to the complement of its previous value. Draw a schematic for a T flip-flop using a D flip-flop and an inverter.

**Exercise 3.7** A *JK flip-flop* receives a clock and two inputs, *J* and *K*. On the rising edge of the clock, it updates the output, *Q*. If *J* and *K* are both 0, *Q* retains its old value. If only *J* is 1, *Q* becomes 1. If only *K* is 1, *Q* becomes 0. If both *J* and *K* are 1, *Q* becomes the opposite of its present state.

- Construct a JK flip-flop using a D flip-flop and some combinational logic.
- Construct a D flip-flop using a JK flip-flop and some combinational logic.
- Construct a T flip-flop (see Exercise 3.6) using a JK flip-flop.

**Exercise 3.8** The circuit in Figure 3.63 is called a *Muller C-element*. Explain in a simple fashion what the relationship is between the inputs and output.

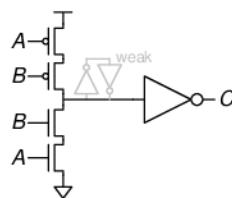


Figure 3.63 Muller C-element

**Exercise 3.9** Design an asynchronously resettable D latch using logic gates.

**Exercise 3.10** Design an asynchronously resettable D flip-flop using logic gates.

**Exercise 3.11** Design a synchronously settable D flip-flop using logic gates.

**Exercise 3.12** Design an asynchronously settable D flip-flop using logic gates.

**Exercise 3.13** Suppose a ring oscillator is built from  $N$  inverters connected in a loop. Each inverter has a minimum delay of  $t_{cd}$  and a maximum delay of  $t_{pd}$ . If  $N$  is odd, determine the range of frequencies at which the oscillator might operate.

**Exercise 3.14** Why must  $N$  be odd in Exercise 3.13?

**Exercise 3.15** Which of the circuits in Figure 3.64 are synchronous sequential circuits? Explain.

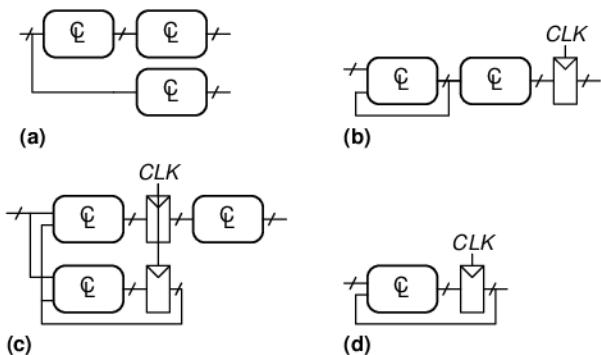


Figure 3.64 Circuits

**Exercise 3.16** You are designing an elevator controller for a building with 25 floors. The controller has two inputs: *UP* and *DOWN*. It produces an output indicating the floor that the elevator is on. There is no floor 13. What is the minimum number of bits of state in the controller?

**Exercise 3.17** You are designing an FSM to keep track of the mood of four students working in the digital design lab. Each student's mood is either **HAPPY** (the circuit works), **SAD** (the circuit blew up), **BUSY** (working on the circuit), **CLUELESS** (confused about the circuit), or **ASLEEP** (face down on the circuit board). How many states does the FSM have? What is the minimum number of bits necessary to represent these states?

**Exercise 3.18** How would you factor the FSM from Exercise 3.17 into multiple simpler machines? How many states does each simpler machine have? What is the minimum total number of bits necessary in this factored design?

**Exercise 3.19** Describe in words what the state machine in Figure 3.65 does.

Using binary state encodings, complete a state transition table and output table for the FSM. Write Boolean equations for the next state and output and sketch a schematic of the FSM.

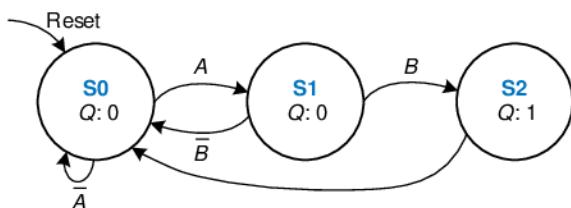


Figure 3.65 State transition diagram

**Exercise 3.20** Describe in words what the state machine in Figure 3.66 does.

Using binary state encodings, complete a state transition table and output table for the FSM. Write Boolean equations for the next state and output and sketch a schematic of the FSM.

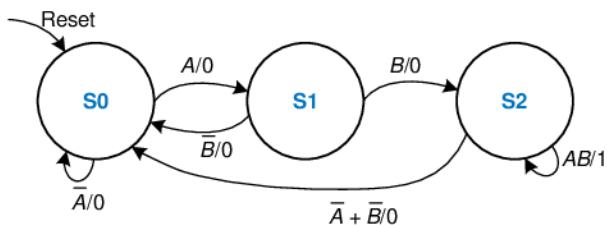


Figure 3.66 State transition diagram

**Exercise 3.21** Accidents are still occurring at the intersection of Academic Avenue and Bravado Boulevard. The football team is rushing into the intersection the moment light B turns green. They are colliding with sleep-deprived CS majors who stagger into the intersection just before light A turns red. Extend the traffic light controller from Section 3.4.1 so that both lights are red for 5 seconds before either light turns green again. Sketch your improved Moore machine state transition diagram, state encodings, state transition table, output table, next state and output equations, and your FSM schematic.

**Exercise 3.22** Alyssa P. Hacker's snail from Section 3.4.3 has a daughter with a Mealy machine FSM brain. The daughter snail smiles whenever she slides over the pattern 1101 or the pattern 1110. Sketch the state transition diagram for this happy snail using as few states as possible. Choose state encodings and write a

combined state transition and output table using your encodings. Write the next state and output equations and sketch your FSM schematic.

**Exercise 3.23** You have been enlisted to design a soda machine dispenser for your department lounge. Sodas are partially subsidized by the student chapter of the IEEE, so they cost only 25 cents. The machine accepts nickels, dimes, and quarters. When enough coins have been inserted, it dispenses the soda and returns any necessary change. Design an FSM controller for the soda machine. The FSM inputs are *Nickel*, *Dime*, and *Quarter*, indicating which coin was inserted. Assume that exactly one coin is inserted on each cycle. The outputs are *Dispense*, *ReturnNickel*, *ReturnDime*, and *ReturnTwoDimes*. When the FSM reaches 25 cents, it asserts *Dispense* and the necessary *Return* outputs required to deliver the appropriate change. Then it should be ready to start accepting coins for another soda.

**Exercise 3.24** Gray codes have a useful property in that consecutive numbers differ in only a single bit position. Table 3.17 lists a 3-bit Gray code representing the numbers 0 to 7. Design a 3-bit modulo 8 Gray code counter FSM with no inputs and three outputs. (A modulo  $N$  counter counts from 0 to  $N - 1$ , then repeats. For example, a watch uses a modulo 60 counter for the minutes and seconds that counts from 0 to 59.) When reset, the output should be 000. On each clock edge, the output should advance to the next Gray code. After reaching 100, it should repeat with 000.

**Table 3.17** 3-bit Gray code

Number	Gray code		
0	0	0	0
1	0	0	1
2	0	1	1
3	0	1	0
4	1	1	0
5	1	1	1
6	1	0	1
7	1	0	0

**Exercise 3.25** Extend your modulo 8 Gray code counter from Exercise 3.24 to be an UP/DOWN counter by adding an *UP* input. If *UP* = 1, the counter advances to the next number. If *UP* = 0, the counter retreats to the previous number.

**Exercise 3.26** Your company, Detect-o-rama, would like to design an FSM that takes two inputs,  $A$  and  $B$ , and generates one output,  $Z$ . The output in cycle  $n$ ,  $Z_n$ , is either the Boolean AND or OR of the corresponding input  $A_n$  and the previous input  $A_{n-1}$ , depending on the other input,  $B_n$ :

$$\begin{aligned} Z_n &= A_n A_{n-1} \quad \text{if } B_n = 0 \\ Z_n &= A_n + A_{n-1} \quad \text{if } B_n = 1 \end{aligned}$$

- (a) Sketch the waveform for  $Z$  given the inputs shown in Figure 3.67.
- (b) Is this FSM a Moore or a Mealy machine?
- (c) Design the FSM. Show your state transition diagram, encoded state transition table, next state and output equations, and schematic.

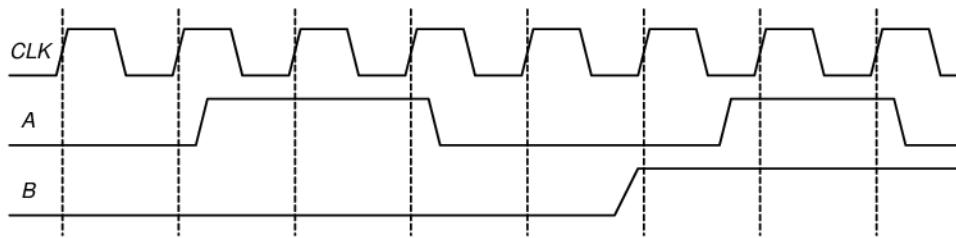


Figure 3.67 FSM input waveforms

**Exercise 3.27** Design an FSM with one input,  $A$ , and two outputs,  $X$  and  $Y$ .  $X$  should be 1 if  $A$  has been 1 for at least three cycles altogether (not necessarily consecutively).  $Y$  should be 1 if  $A$  has been 1 for at least two consecutive cycles. Show your state transition diagram, encoded state transition table, next state and output equations, and schematic.

**Exercise 3.28** Analyze the FSM shown in Figure 3.68. Write the state transition and output tables and sketch the state transition diagram. Describe in words what the FSM does.

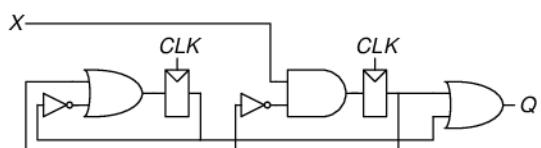


Figure 3.68 FSM schematic

**Exercise 3.29** Repeat Exercise 3.28 for the FSM shown in Figure 3.69. Recall that the  $r$  and  $s$  register inputs indicate set and reset, respectively.

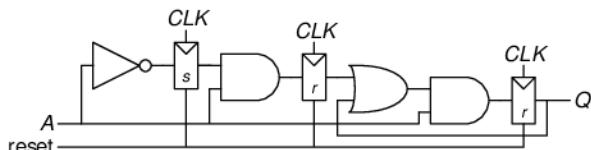


Figure 3.69 FSM schematic

**Exercise 3.30** Ben Bitdiddle has designed the circuit in Figure 3.70 to compute a registered four-input XOR function. Each two-input XOR gate has a propagation delay of 100 ps and a contamination delay of 55 ps. Each flip-flop has a setup time of 60 ps, a hold time of 20 ps, a clock-to- $Q$  maximum delay of 70 ps, and a clock-to- $Q$  minimum delay of 50 ps.

- (a) If there is no clock skew, what is the maximum operating frequency of the circuit?
- (b) How much clock skew can the circuit tolerate if it must operate at 2 GHz?
- (c) How much clock skew can the circuit tolerate before it might experience a hold time violation?
- (d) Alyssa P. Hacker points out that she can redesign the combinational logic between the registers to be faster *and* tolerate more clock skew. Her improved circuit also uses three two-input XORs, but they are arranged differently. What is her circuit? What is its maximum frequency if there is no clock skew? How much clock skew can the circuit tolerate before it might experience a hold time violation?

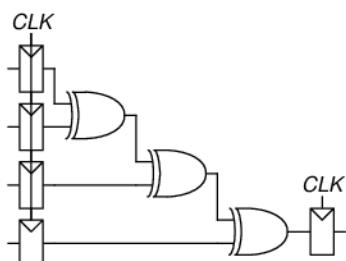
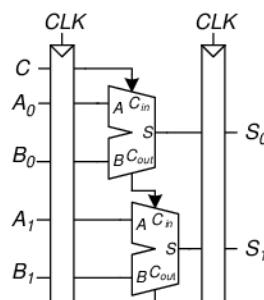


Figure 3.70 Registered four-input XOR circuit

**Exercise 3.31** You are designing an adder for the blindingly fast 2-bit RePentium Processor. The adder is built from two full adders such that the carry out of the first adder is the carry in to the second adder, as shown in Figure 3.71. Your adder has input and output registers and must complete the



**Figure 3.71** 2-bit adder schematic

addition in one clock cycle. Each full adder has the following propagation delays: 20 ps from  $C_{in}$  to  $C_{out}$  or to *Sum* ( $S$ ), 25 ps from  $A$  or  $B$  to  $C_{out}$ , and 30 ps from  $A$  or  $B$  to  $S$ . The adder has a contamination delay of 15 ps from  $C_{in}$  to either output and 22 ps from  $A$  or  $B$  to either output. Each flip-flop has a setup time of 30 ps, a hold time of 10 ps, a clock-to- $Q$  propagation delay of 35 ps, and a clock-to- $Q$  contamination delay of 21 ps.

- (a) If there is no clock skew, what is the maximum operating frequency of the circuit?
- (b) How much clock skew can the circuit tolerate if it must operate at 8 GHz?
- (c) How much clock skew can the circuit tolerate before it might experience a hold time violation?

**Exercise 3.32** A field programmable gate array (FPGA) uses *configurable logic blocks* (CLBs) rather than logic gates to implement combinational logic. The Xilinx Spartan 3 FPGA has propagation and contamination delays of 0.61 and 0.30 ns, respectively for each CLB. It also contains flip-flops with propagation and contamination delays of 0.72 and 0.50 ns, and setup and hold times of 0.53 and 0 ns, respectively.

- (a) If you are building a system that needs to run at 40 MHz, how many consecutive CLBs can you use between two flip-flops? Assume there is no clock skew and no delay through wires between CLBs.
- (b) Suppose that all paths between flip-flops pass through at least one CLB. How much clock skew can the FPGA have without violating the hold time?

**Exercise 3.33** A synchronizer is built from a pair of flip-flops with  $t_{\text{setup}} = 50$  ps,  $T_0 = 20$  ps, and  $\tau = 30$  ps. It samples an asynchronous input that changes  $10^8$  times per second. What is the minimum clock period of the synchronizer to achieve a mean time between failures (MTBF) of 100 years?

**Exercise 3.34** You would like to build a synchronizer that can receive asynchronous inputs with an MTBF of 50 years. Your system is running at 1 GHz, and you use sampling flip-flops with  $\tau = 100$  ps,  $T_0 = 110$  ps, and  $t_{\text{setup}} = 70$  ps. The synchronizer receives a new asynchronous input on average 0.5 times per second (i.e., once every 2 seconds). What is the required probability of failure to satisfy this MTBF? How many clock cycles would you have to wait before reading the sampled input signal to give that probability of error?

**Exercise 3.35** You are walking down the hallway when you run into your lab partner walking in the other direction. The two of you first step one way and are still in each other's way. Then you both step the other way and are still in each other's way. Then you both wait a bit, hoping the other person will step aside. You can model this situation as a metastable point and apply the same theory that has been applied to synchronizers and flip-flops. Suppose you create a mathematical model for yourself and your lab partner. You start the unfortunate encounter in the metastable state. The probability that you remain in this state after  $t$  seconds is  $e^{-\frac{t}{\tau}}$ .  $\tau$  indicates your response rate; today, your brain has been blurred by lack of sleep and has  $\tau = 20$  seconds.

- (a) How long will it be until you have 99% certainty that you will have resolved from metastability (i.e., figured out how to pass one another)?
- (b) You are not only sleepy, but also ravenously hungry. In fact, you will starve to death if you don't get going to the cafeteria within 3 minutes. What is the probability that your lab partner will have to drag you to the morgue?

**Exercise 3.36** You have built a synchronizer using flip-flops with  $T_0 = 20$  ps and  $\tau = 30$  ps. Your boss tells you that you need to increase the MTBF by a factor of 10. By how much do you need to increase the clock period?

**Exercise 3.37** Ben Bitdiddle invents a new and improved synchronizer in Figure 3.72 that he claims eliminates metastability in a single cycle. He explains that the circuit in box  $M$  is an analog “metastability detector” that produces a HIGH output if the input voltage is in the forbidden zone between  $V_{IL}$  and  $V_{IH}$ . The metastability detector checks to determine

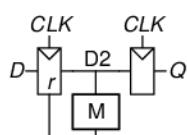


Figure 3.72 “New and improved” synchronizer

whether the first flip-flop has produced a metastable output on  $D_2$ . If so, it asynchronously resets the flip-flop to produce a good 0 at  $D_2$ . The second flip-flop then samples  $D_2$ , always producing a valid logic level on  $Q$ . Alyssa P. Hacker tells Ben that there must be a bug in the circuit, because eliminating metastability is just as impossible as building a perpetual motion machine. Who is right? Explain, showing Ben's error or showing why Alyssa is wrong.

## Interview Questions

The following exercises present questions that have been asked at interviews for digital design jobs.

**Question 3.1** Draw a state machine that can detect when it has received the serial input sequence 01010.

**Question 3.2** Design a serial (one bit at a time) two's complementer FSM with two inputs, *Start* and *A*, and one output, *Q*. A binary number of arbitrary length is provided to input *A*, starting with the least significant bit. The corresponding bit of the output appears at *Q* on the same cycle. *Start* is asserted for one cycle to initialize the FSM before the least significant bit is provided.

**Question 3.3** What is the difference between a latch and a flip-flop? Under what circumstances is each one preferable?

**Question 3.4** Design a 5-bit counter finite state machine.

**Question 3.5** Design an edge detector circuit. The output should go HIGH for one cycle after the input makes a  $0 \rightarrow 1$  transition.

**Question 3.6** Describe the concept of pipelining and why it is used.

**Question 3.7** Describe what it means for a flip-flop to have a negative hold time.

**Question 3.8** Given signal *A*, shown in Figure 3.73, design a circuit that produces signal *B*.

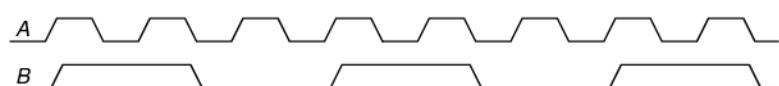


Figure 3.73 Signal waveforms

**Question 3.9** Consider a block of logic between two registers. Explain the timing constraints. If you add a buffer on the clock input of the receiver (the second flip-flop), does the setup time constraint get better or worse?



# 4

## Hardware Description Languages

4.1	<a href="#">Introduction</a>
4.2	<a href="#">Combinational Logic</a>
4.3	<a href="#">Structural Modeling</a>
4.4	<a href="#">Sequential Logic</a>
4.5	<a href="#">More Combinational Logic</a>
4.6	<a href="#">Finite State Machines</a>
4.7	<a href="#">Parameterized Modules*</a>
4.8	<a href="#">Testbenches</a>
4.9	<a href="#">Summary</a>
	<a href="#">Exercises</a>
	<a href="#">Interview Questions</a>

### 4.1 INTRODUCTION

Thus far, we have focused on designing combinational and sequential digital circuits at the schematic level. The process of finding an efficient set of logic gates to perform a given function is labor intensive and error prone, requiring manual simplification of truth tables or Boolean equations and manual translation of finite state machines (FSMs) into gates. In the 1990's, designers discovered that they were far more productive if they worked at a higher level of abstraction, specifying just the logical function and allowing a *computer-aided design* (CAD) tool to produce the optimized gates. The specifications are generally given in a *hardware description language* (HDL). The two leading hardware description languages are *Verilog* and *VHDL*.

Verilog and VHDL are built on similar principles but have different syntax. Discussion of these languages in this chapter is divided into two columns for literal side-by-side comparison, with Verilog on the left and VHDL on the right. When you read the chapter for the first time, focus on one language or the other. Once you know one, you'll quickly master the other if you need it.

Subsequent chapters show hardware in both schematic and HDL form. If you choose to skip this chapter and not learn one of the HDLs, you will still be able to master the principles of computer organization from the schematics. However, the vast majority of commercial systems are now built using HDLs rather than schematics. If you expect to do digital design at any point in your professional life, we urge you to learn one of the HDLs.

#### 4.1.1 Modules

A block of hardware with inputs and outputs is called a *module*. An AND gate, a multiplexer, and a priority circuit are all examples of hardware modules. The two general styles for describing module functionality are

*behavioral* and *structural*. Behavioral models describe what a module does. Structural models describe how a module is built from simpler pieces; it is an application of hierarchy. The Verilog and VHDL code in HDL Example 4.1 illustrate behavioral descriptions of a module that computes the Boolean function from Example 2.6,  $y = \overline{a}\overline{b}\overline{c} + a\overline{b}\overline{c} + a\overline{b}c$ . In both languages, the module is named `sillyfunction` and has three inputs, `a`, `b`, and `c`, and one output, `y`.

#### HDL Example 4.1 COMBINATIONAL LOGIC

##### Verilog

```
module sillyfunction (input a, b, c,
                      output y);

    assign y = ~a & ~b & ~c |
               a & ~b & ~c |
               a & ~b & c;

endmodule
```

A Verilog module begins with the module name and a listing of the inputs and outputs. The `assign` statement describes combinational logic. `~` indicates NOT, `&` indicates AND, and `|` indicates OR.

Verilog signals such as the inputs and outputs are Boolean variables (0 or 1). They may also have floating and undefined values, as discussed in Section 4.2.8.

##### VHDL

```
library IEEE; use IEEE.STD_LOGIC_1164.all;

entity sillyfunction is
    port (a, b, c: in STD_LOGIC;
          y:        out STD_LOGIC);
end;

architecture synth of sillyfunction is
begin
    y <= ((not a) and (not b) and (not c)) or
         (a and (not b) and (not c)) or
         (a and (not b) and c);
end;
```

VHDL code has three parts: the library use clause, the entity declaration, and the architecture body. The library use clause is required and will be discussed in Section 4.2.11. The entity declaration lists the module name and its inputs and outputs. The architecture body defines what the module does.

VHDL signals, such as inputs and outputs, must have a *type declaration*. Digital signals should be declared to be `STD_LOGIC` type. `STD_LOGIC` signals can have a value of '0' or '1', as well as floating and undefined values that will be described in Section 4.2.8. The `STD_LOGIC` type is defined in the `IEEE.STD_LOGIC_1164` library, which is why the library must be used.

VHDL lacks a good default order of operations, so Boolean equations should be parenthesized.

A module, as you might expect, is a good application of modularity. It has a well defined interface, consisting of its inputs and outputs, and it performs a specific function. The particular way in which it is coded is unimportant to others that might use the module, as long as it performs its function.

#### 4.1.2 Language Origins

Universities are almost evenly split on which of these languages is taught in a first course, and industry is similarly split on which language is preferred. Compared to Verilog, VHDL is more verbose and cumbersome,

### Verilog

Verilog was developed by Gateway Design Automation as a proprietary language for logic simulation in 1984. Gateway was acquired by Cadence in 1989 and Verilog was made an open standard in 1990 under the control of Open Verilog International. The language became an IEEE standard<sup>1</sup> in 1995 (IEEE STD 1364) and was updated in 2001.

### VHDL

VHDL is an acronym for the *VHSIC Hardware Description Language*. VHSIC is in turn an acronym for the *Very High Speed Integrated Circuits* program of the US Department of Defense.

VHDL was originally developed in 1981 by the Department of Defense to describe the structure and function of hardware. Its roots draw from the Ada programming language. The IEEE standardized it in 1987 (IEEE STD 1076) and has updated the standard several times since. The language was first envisioned for documentation but was quickly adopted for simulation and synthesis.

as you might expect of a language developed by committee. U.S. military contractors, the European Space Agency, and telecommunications companies use VHDL extensively.

Both languages are fully capable of describing any hardware system, and both have their quirks. The best language to use is the one that is already being used at your site or the one that your customers demand. Most CAD tools today allow the two languages to be mixed, so that different modules can be described in different languages.

#### 4.1.3 Simulation and Synthesis

The two major purposes of HDLs are logic *simulation* and *synthesis*. During simulation, inputs are applied to a module, and the outputs are checked to verify that the module operates correctly. During synthesis, the textual description of a module is transformed into logic gates.

#### Simulation

Humans routinely make mistakes. Such errors in hardware designs are called *bugs*. Eliminating the bugs from a digital system is obviously important, especially when customers are paying money and lives depend on the correct operation. Testing a system in the laboratory is time-consuming. Discovering the cause of errors in the lab can be extremely difficult, because only signals routed to the chip pins can be observed. There is no way to directly observe what is happening inside a chip. Correcting errors after the system is built can be devastatingly expensive. For example, correcting a mistake in a cutting-edge integrated circuit costs more than a million dollars and takes several months. Intel's infamous FDIV (floating point division) bug in the Pentium processor forced the company to recall chips after they had shipped, at a total cost of \$475 million. Logic simulation is essential to test a system before it is built.

The term “bug” predates the invention of the computer. Thomas Edison called the “little faults and difficulties” with his inventions “bugs” in 1878.

The first real computer bug was a moth, which got caught between the relays of the Harvard Mark II electro-mechanical computer in 1947. It was found by Grace Hopper, who logged the incident, along with the moth itself and the comment “first actual case of bug being found.”



Source: Notebook entry courtesy Naval Historical Center, US Navy; photo No. NII 96566-KN

<sup>1</sup> The Institute of Electrical and Electronics Engineers (IEEE) is a professional society responsible for many computing standards including WiFi (802.11), Ethernet (802.3), and floating-point numbers (754) (see Chapter 5).

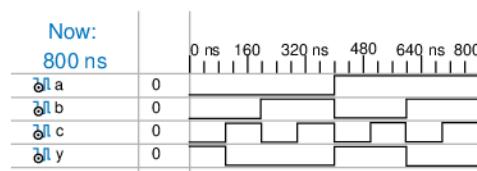
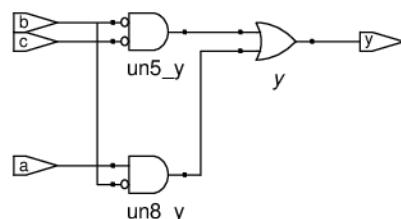
**Figure 4.1** Simulation waveforms

Figure 4.1 shows waveforms from a simulation<sup>2</sup> of the previous `sillyfunction` module demonstrating that the module works correctly. `y` is TRUE when `a`, `b`, and `c` are 000, 100, or 101, as specified by the Boolean equation.

### Synthesis

Logic synthesis transforms HDL code into a *netlist* describing the hardware (e.g., the logic gates and the wires connecting them). The logic synthesizer might perform optimizations to reduce the amount of hardware required. The netlist may be a text file, or it may be drawn as a schematic to help visualize the circuit. Figure 4.2 shows the results of synthesizing the `sillyfunction` module.<sup>3</sup> Notice how the three three-input AND gates are simplified into two two-input AND gates, as we discovered in Example 2.6 using Boolean algebra.

Circuit descriptions in HDL resemble code in a programming language. However, you must remember that the code is intended to represent hardware. Verilog and VHDL are rich languages with many commands. Not all of these commands can be synthesized into hardware. For example, a command to print results on the screen during simulation does not translate into hardware. Because our primary interest is

**Figure 4.2** Synthesized circuit

<sup>2</sup> The simulation was performed with the Xilinx ISE Simulator, which is part of the Xilinx ISE 8.2 software. The simulator was selected because it is used commercially, yet is freely available to universities.

<sup>3</sup> Synthesis was performed with Synplify Pro from Synplicity. The tool was selected because it is the leading commercial tool for synthesizing HDL to field-programmable gate arrays (see Section 5.6.2) and because it is available inexpensively for universities.

to build hardware, we will emphasize a *synthesizable subset* of the languages. Specifically, we will divide HDL code into *synthesizable* modules and a *testbench*. The synthesizable modules describe the hardware. The testbench contains code to apply inputs to a module, check whether the output results are correct, and print discrepancies between expected and actual outputs. Testbench code is intended only for simulation and cannot be synthesized.

One of the most common mistakes for beginners is to think of HDL as a computer program rather than as a shorthand for describing digital hardware. If you don't know approximately what hardware your HDL should synthesize into, you probably won't like what you get. You might create far more hardware than is necessary, or you might write code that simulates correctly but cannot be implemented in hardware. Instead, think of your system in terms of blocks of combinational logic, registers, and finite state machines. Sketch these blocks on paper and show how they are connected before you start writing code.

In our experience, the best way to learn an HDL is by example. HDLs have specific ways of describing various classes of logic; these ways are called *idioms*. This chapter will teach you how to write the proper HDL idioms for each type of block and then how to put the blocks together to produce a working system. When you need to describe a particular kind of hardware, look for a similar example and adapt it to your purpose. We do not attempt to rigorously define all the syntax of the HDLs, because that is deathly boring and because it tends to encourage thinking of HDLs as programming languages, not shorthand for hardware. The IEEE Verilog and VHDL specifications, and numerous dry but exhaustive textbooks, contain all of the details, should you find yourself needing more information on a particular topic. (See Further Readings section at back of the book.)

## 4.2 COMBINATIONAL LOGIC

Recall that we are disciplining ourselves to design synchronous sequential circuits, which consist of combinational logic and registers. The outputs of combinational logic depend only on the current inputs. This section describes how to write behavioral models of combinational logic with HDLs.

### 4.2.1 Bitwise Operators

*Bitwise* operators act on single-bit signals or on multi-bit busses. For example, the `inv` module in HDL Example 4.2 describes four inverters connected to 4-bit busses.

**HDL Example 4.2 INVERTERS****Verilog**

```
module inv (input [3:0] a,
            output [3:0] y);

    assign y = ~a;
endmodule
```

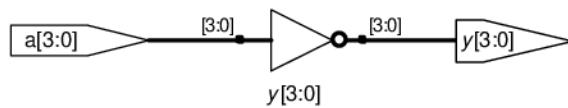
`a[3:0]` represents a 4-bit bus. The bits, from most significant to least significant, are `a[3]`, `a[2]`, `a[1]`, and `a[0]`. This is called *little-endian* order, because the least significant bit has the smallest bit number. We could have named the bus `a[4:1]`, in which case `a[4]` would have been the most significant. Or we could have used `a[0:3]`, in which case the bits, from most significant to least significant, would be `a[0]`, `a[1]`, `a[2]`, and `a[3]`. This is called *big-endian* order.

**VHDL**

```
library IEEE; use IEEE.STD_LOGIC_1164.all;
entity inv is
    port (a: in STD_LOGIC_VECTOR (3 downto 0);
          y: out STD_LOGIC_VECTOR (3 downto 0));
end;
```

```
architecture synth of inv is
begin
    y <= not a;
end;
```

VHDL uses `STD_LOGIC_VECTOR`, to indicate busses of `STD_LOGIC`. `STD_LOGIC_VECTOR (3 downto 0)` represents a 4-bit bus. The bits, from most significant to least significant, are 3, 2, 1, and 0. This is called *little-endian* order, because the least significant bit has the smallest bit number. We could have declared the bus to be `STD_LOGIC_VECTOR (4 downto 1)`, in which case bit 4 would have been the most significant. Or we could have written `STD_LOGIC_VECTOR (0 to 3)`, in which case the bits, from most significant to least significant, would be 0, 1, 2, and 3. This is called *big-endian* order.



**Figure 4.3** `inv` synthesized circuit

The endianness of a bus is purely arbitrary. (See the sidebar in Section 6.2.2 for the origin of the term.) Indeed, endianness is also irrelevant to this example, because a bank of inverters doesn't care what the order of the bits are. Endianness matters only for operators, such as addition, where the sum of one column carries over into the next. Either ordering is acceptable, as long as it is used consistently. We will consistently use the little-endian order, `[N-1:0]` in Verilog and `(N-1 downto 0)` in VHDL, for an  $N$ -bit bus.

After each code example in this chapter is a schematic produced from the Verilog code by the Synplify Pro synthesis tool. Figure 4.3 shows that the `inv` module synthesizes to a bank of four inverters, indicated by the inverter symbol labeled `y[3:0]`. The bank of inverters connects to 4-bit input and output busses. Similar hardware is produced from the synthesized VHDL code.

The gates module in HDL Example 4.3 demonstrates bitwise operations acting on 4-bit busses for other basic logic functions.

**HDL Example 4.3 LOGIC GATES****Verilog**

```
module gates (input [3:0] a, b,
              output [3:0] y1, y2,
              y3, y4, y5);

  /* Five different two-input logic
   gates acting on 4 bit busses */
  assign y1 = a & b; // AND
  assign y2 = a | b; // OR
  assign y3 = a ^ b; // XOR
  assign y4 = ~(a & b); // NAND
  assign y5 = ~(a | b); // NOR
endmodule
```

$\sim$ ,  $\wedge$ , and  $\mid$  are examples of Verilog *operators*, whereas  $a$ ,  $b$ , and  $y_1$  are *operands*. A combination of operators and operands, such as  $a \wedge b$ , or  $\sim(a \mid b)$ , is called an *expression*. A complete command such as `assign y4 = ~(a & b);` is called a *statement*.

`assign out = in1 op in2;` is called a *continuous assignment statement*. Continuous assignment statements end with a semicolon. Anytime the inputs on the right side of the `=` in a continuous assignment statement change, the output on the left side is recomputed. Thus, continuous assignment statements describe combinational logic.

**VHDL**

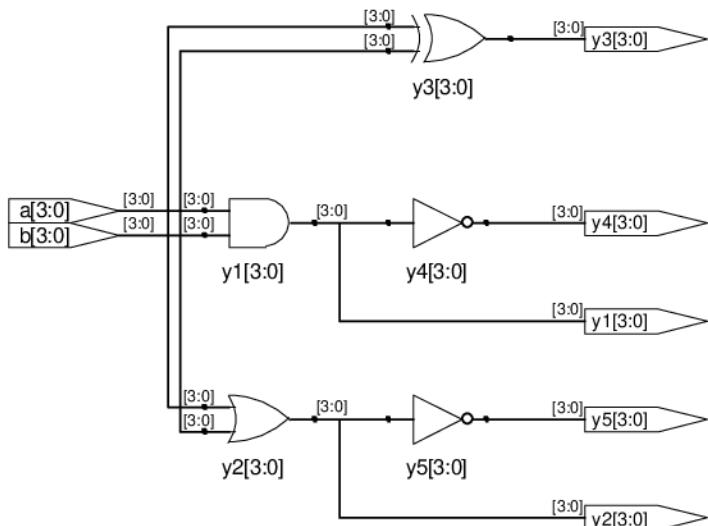
```
library IEEE; use IEEE.STD_LOGIC_1164.all;

entity gates is
  port(a, b: in STD_LOGIC_VECTOR (3 downto 0);
       y1, y2, y3, y4,
       y5: out STD_LOGIC_VECTOR (3 downto 0));
end;

architecture synth of gates is
begin
  -- Five different two-input logic gates
  -- acting on 4 bit busses
  y1 <= a and b;
  y2 <= a or b;
  y3 <= a xor b;
  y4 <= a nand b;
  y5 <= a nor b;
end;
```

`not`, `xor`, `and` or `or` are examples of VHDL *operators*, whereas  $a$ ,  $b$ , and  $y_1$  are *operands*. A combination of operators and operands, such as  $a$  and  $b$ , or  $a$  nor  $b$ , is called an *expression*. A complete command such as `y4 <= a nand b;` is called a *statement*.

`out <= in1 op in2;` is called a *concurrent signal assignment statement*. VHDL assignment statements end with a semicolon. Anytime the inputs on the right side of the `<=` in a concurrent signal assignment statement change, the output on the left side is recomputed. Thus, concurrent signal assignment statements describe combinational logic.



**Figure 4.4** gates synthesized circuit

### 4.2.2 Comments and White Space

The gates example showed how to format comments. Verilog and VHDL are not picky about the use of white space (i.e., spaces, tabs, and line breaks). Nevertheless, proper indenting and use of blank lines is helpful to make nontrivial designs readable. Be consistent in your use of capitalization and underscores in signal and module names. Module and signal names must not begin with a digit.

#### Verilog

Verilog comments are just like those in C or Java. Comments beginning with `/*` continue, possibly across multiple lines, to the next `*/`. Comments beginning with `//` continue to the end of the line.

Verilog is case-sensitive. `y1` and `Y1` are different signals in Verilog.

#### VHDL

VHDL comments begin with `--` and continue to the end of the line. Comments spanning multiple lines must use `--` at the beginning of each line.

VHDL is not case-sensitive. `y1` and `Y1` are the same signal in VHDL. However, other tools that may read your file might be case sensitive, leading to nasty bugs if you blithely mix upper and lower case.

### 4.2.3 Reduction Operators

Reduction operators imply a multiple-input gate acting on a single bus. HDL Example 4.4 describes an eight-input AND gate with inputs  $a_7, a_6, \dots, a_0$ .

---

#### HDL Example 4.4 EIGHT-INPUT AND

#### Verilog

```
module and8(input [7:0] a,
             output      y);

  assign y = &a;

  // &a is much easier to write than
  // assign y = a[7] & a[6] & a[5] & a[4] &
  //           a[3] & a[2] & a[1] & a[0];
endmodule
```

As one would expect, `|`, `^`, `~&`, and `~ |` reduction operators are available for OR, XOR, NAND, and NOR as well. Recall that a multi-input XOR performs parity, returning TRUE if an odd number of inputs are TRUE.

#### VHDL

VHDL does not have reduction operators. Instead, it provides the `generate` command (see Section 4.7). Alternatively, the operation can be written explicitly, as shown below.

```
library IEEE; use IEEE.STD_LOGIC_1164.all;

entity and8 is
  port (a: in STD_LOGIC_VECTOR(7 downto 0);
        y: out STD_LOGIC);
end;

architecture synth of and8 is
begin
  y <= a(7) and a(6) and a(5) and a(4) and
      a(3) and a(2) and a(1) and a(0);
end;
```

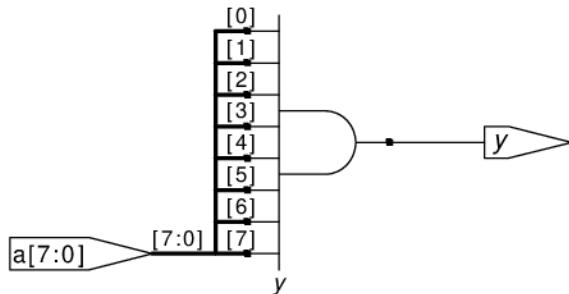


Figure 4.5 and8 synthesized circuit

#### 4.2.4 Conditional Assignment

*Conditional assignments* select the output from among alternatives based on an input called the *condition*. HDL Example 4.5 illustrates a 2:1 multiplexer using conditional assignment.

##### HDL Example 4.5 2:1 MULTIPLEXER

###### Verilog

The *conditional operator* `?:` chooses, based on a first expression, between a second and third expression. The first expression is called the *condition*. If the condition is 1, the operator chooses the second expression. If the condition is 0, the operator chooses the third expression.

`?:` is especially useful for describing a multiplexer because, based on the first input, it selects between two others. The following code demonstrates the idiom for a 2:1 multiplexer with 4-bit inputs and outputs using the conditional operator.

```
module mux2(input [3:0] d0, d1,
             input     s,
             output [3:0] y);

    assign y = s ? d1 : d0;
endmodule
```

If  $s$  is 1, then  $y = d1$ . If  $s$  is 0, then  $y = d0$ .

`?:` is also called a *ternary operator*, because it takes three inputs. It is used for the same purpose in the C and Java programming languages.

###### VHDL

*Conditional signal assignments* perform different operations depending on some condition. They are especially useful for describing a multiplexer. For example, a 2:1 multiplexer can use conditional signal assignment to select one of two 4-bit inputs.

```
library IEEE; use IEEE.STD_LOGIC_1164.all;

entity mux2 is
    port(d0, d1: in STD_LOGIC_VECTOR(3 downto 0);
         s:      in STD_LOGIC;
         y:      out STD_LOGIC_VECTOR(3 downto 0));
end;

architecture synth of mux2 is
begin
    y <= d0 when s = '0' else d1;
end;
```

The conditional signal assignment sets  $y$  to  $d0$  if  $s$  is 0. Otherwise it sets  $y$  to  $d1$ .

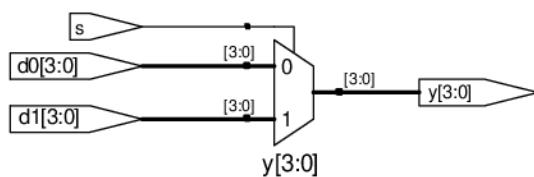


Figure 4.6 mux2 synthesized circuit

HDL Example 4.6 shows a 4:1 multiplexer based on the same principle as the 2:1 multiplexer in HDL Example 4.5.

Figure 4.7 shows the schematic for the 4:1 multiplexer produced by Synplify Pro. The software uses a different multiplexer symbol than this text has shown so far. The multiplexer has multiple data (d) and one-hot enable (e) inputs. When one of the enables is asserted, the associated data is passed to the output. For example, when  $s[1] = s[0] = 0$ , the bottom AND gate, unl\_s\_5, produces a 1, enabling the bottom input of the multiplexer and causing it to select  $d0[3:0]$ .

#### HDL Example 4.6 4:1 MULTIPLEXER

##### Verilog

A 4:1 multiplexer can select one of four inputs using nested conditional operators.

```
module mux4(input [3:0] d0, d1, d2, d3,
             input [1:0] s,
             output [3:0] y);

    assign y = s[1] ? (s[0] ? d3 : d2)
                  : (s[0] ? d1 : d0);
endmodule
```

If  $s[1]$  is 1, then the multiplexer chooses the first expression,  $(s[0] ? d3 : d2)$ . This expression in turn chooses either  $d3$  or  $d2$  based on  $s[0]$  ( $y = d3$  if  $s[0]$  is 1 and  $d2$  if  $s[0]$  is 0). If  $s[1]$  is 0, then the multiplexer similarly chooses the second expression, which gives either  $d1$  or  $d0$  based on  $s[0]$ .

##### VHDL

A 4:1 multiplexer can select one of four inputs using multiple `else` clauses in the conditional signal assignment.

```
library IEEE; use IEEE.STD_LOGIC_1164.all;
entity mux4 is
    port (d0, d1,
          d2, d3: in STD_LOGIC_VECTOR(3 downto 0);
          s:      in STD_LOGIC_VECTOR(1 downto 0);
          y:      out STD_LOGIC_VECTOR(3 downto 0));
end;

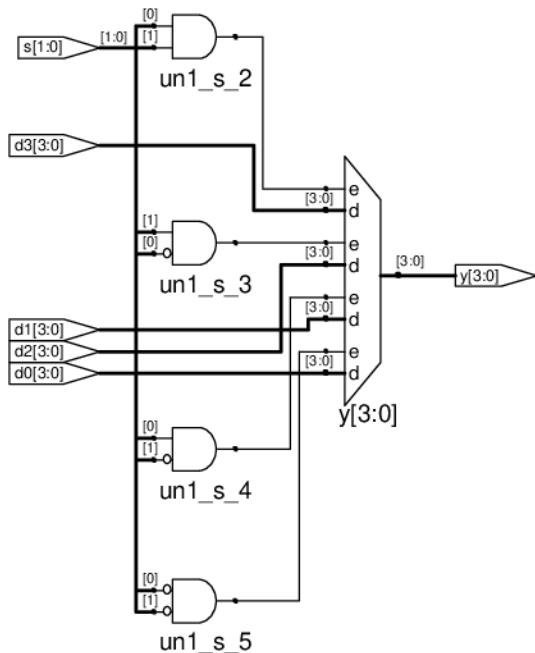
architecture synth1 of mux4 is
begin
    y <= d0 when s = "00" else
              d1 when s = "01" else
              d2 when s = "10" else
              d3;
end;
```

VHDL also supports *selected signal assignment statements* to provide a shorthand when selecting from one of several possibilities. This is analogous to using a `case` statement in place of multiple `if/else` statements in some programming languages. The 4:1 multiplexer can be rewritten with selected signal assignment as follows:

```
architecture synth2 of mux4 is
begin
    with s select y <=
        d0 when "00",
        d1 when "01",
        d2 when "10",
        d3 when others;
end;
```

#### 4.2.5 Internal Variables

Often it is convenient to break a complex function into intermediate steps. For example, a full adder, which will be described in Section 5.2.1,



**Figure 4.7** mux4 synthesized circuit

is a circuit with three inputs and two outputs defined by the following equations:

$$\begin{aligned} S &= A \oplus B \oplus C_{in} \\ C_{out} &= AB + AC_{in} + BC_{in} \end{aligned} \quad (4.1)$$

If we define intermediate signals,  $P$  and  $G$ ,

$$\begin{aligned} P &= A \oplus B \\ G &= AB \end{aligned} \quad (4.2)$$

we can rewrite the full adder as follows:

$$\begin{aligned} S &= P \oplus C_{in} \\ C_{out} &= G + PC_{in} \end{aligned} \quad (4.3)$$

$P$  and  $G$  are called *internal variables*, because they are neither inputs nor outputs but are used only internal to the module. They are similar to local variables in programming languages. HDL Example 4.7 shows how they are used in HDLs.

Check this by filling out the truth table to convince yourself it is correct.

HDL assignment statements (assign in Verilog and  $<=$  in VHDL) take place concurrently. This is different from conventional programming languages such as C or Java, in which statements are evaluated in the order in which they are written. In a conventional language, it is